

Cracks in the Shield: Understanding and Mitigating Side-Channel Threats to Confidential Computing

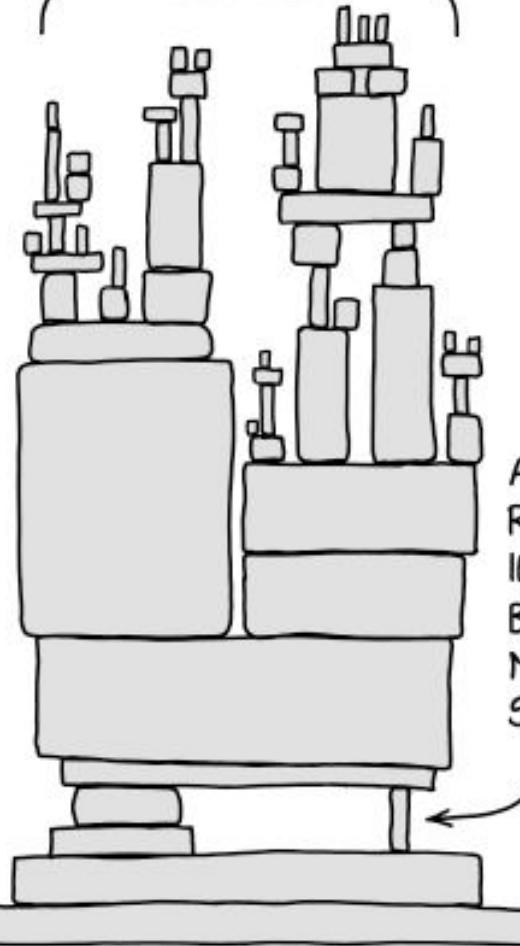
Jo Van Bulck

🏡 DistriNet, KU Leuven, Belgium 📩 jo.vanbulck@cs.kuleuven.be 🌐 vanbulck.net

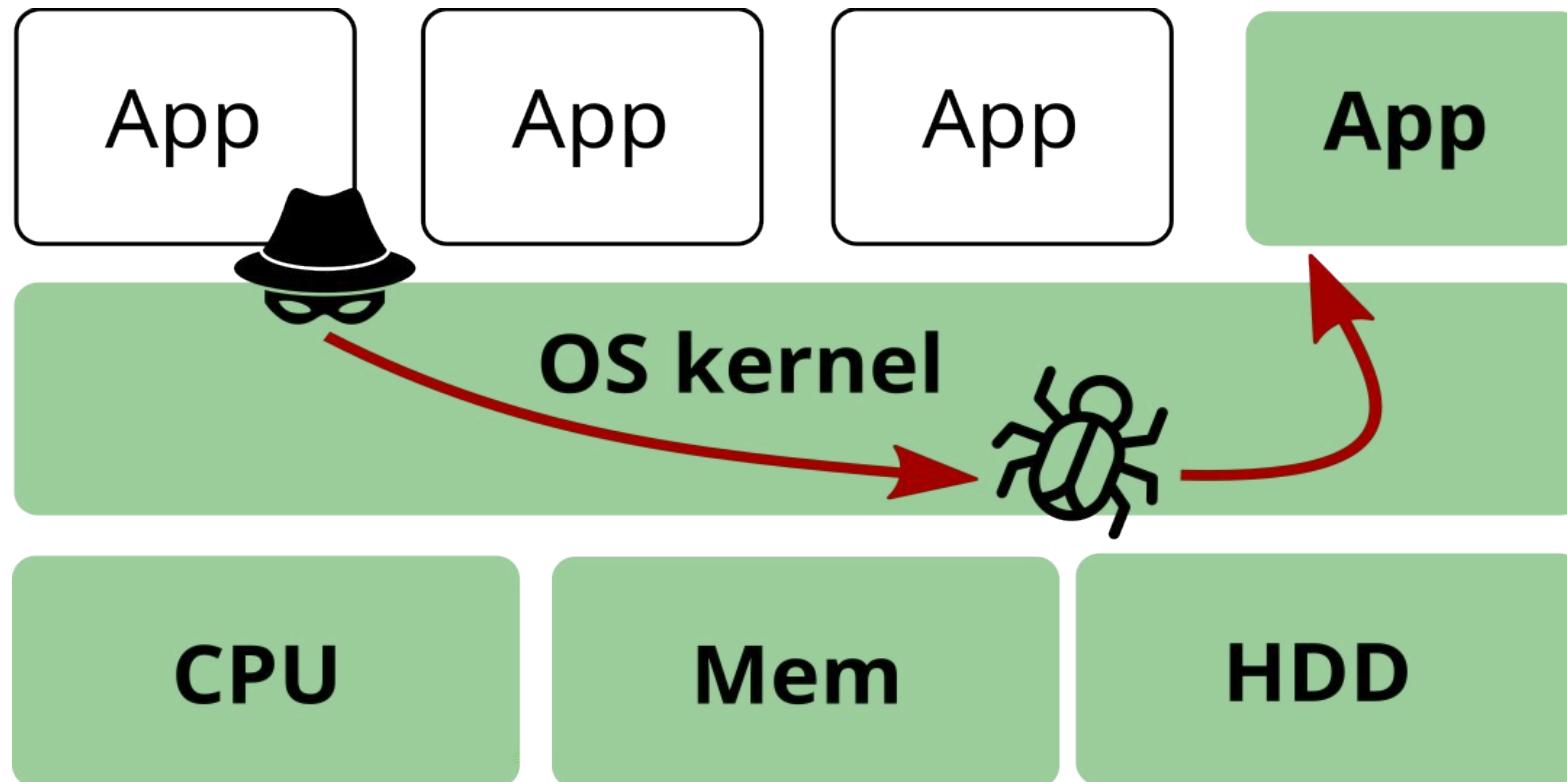
ETH, Sept 22, 2025

Trust?

ALL MODERN DIGITAL
INFRASTRUCTURE

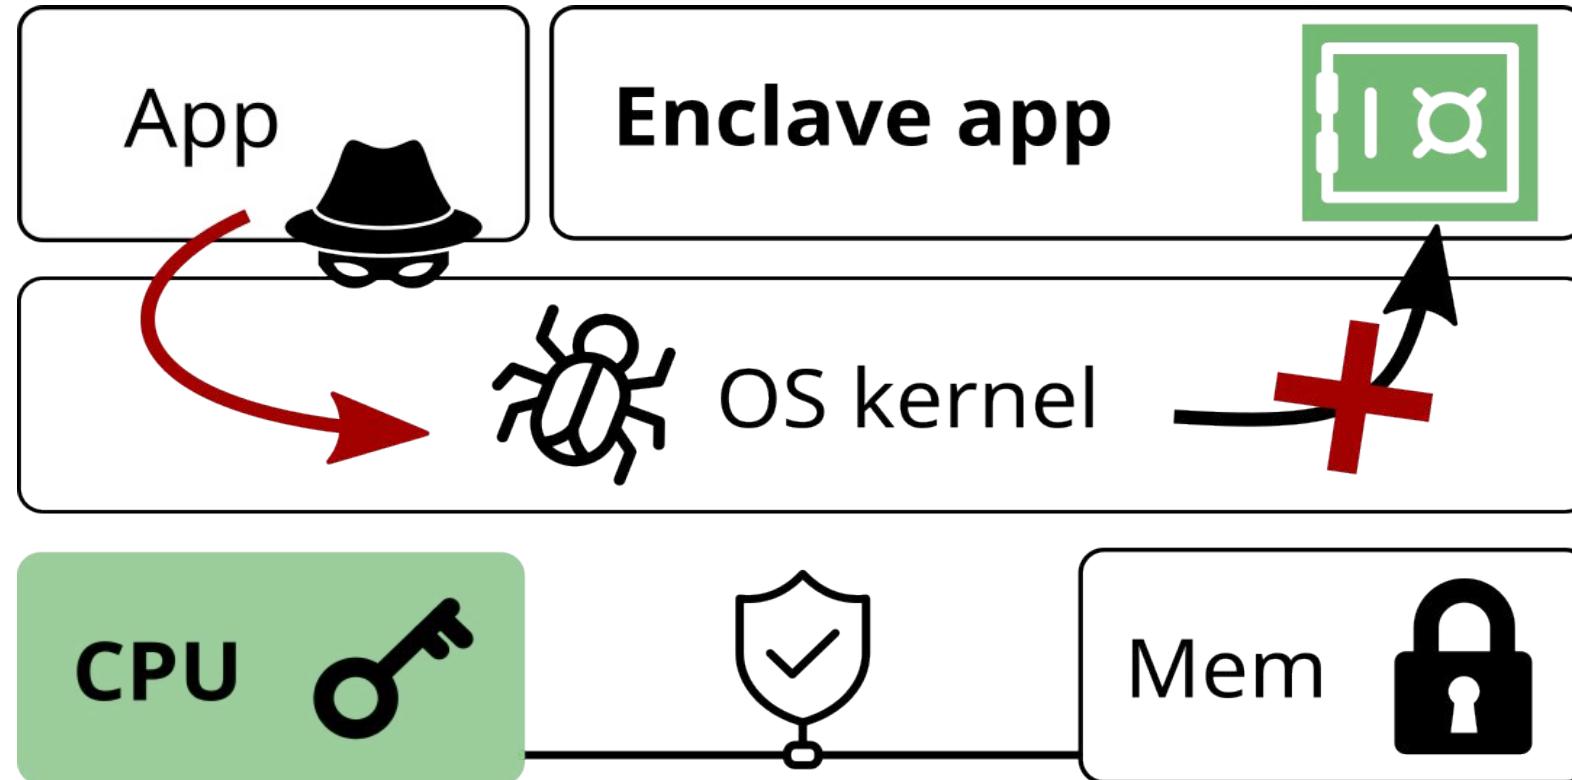


Confidential Computing: Reducing Attack Surface



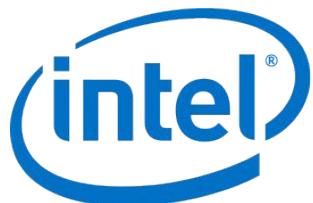
Traditional layered designs: Large **trusted computing base**

Confidential Computing: Reducing Attack Surface



Trusted execution: Hardware-level **isolation** and attestation

The Rise of Trusted Execution Environments (TEEs)

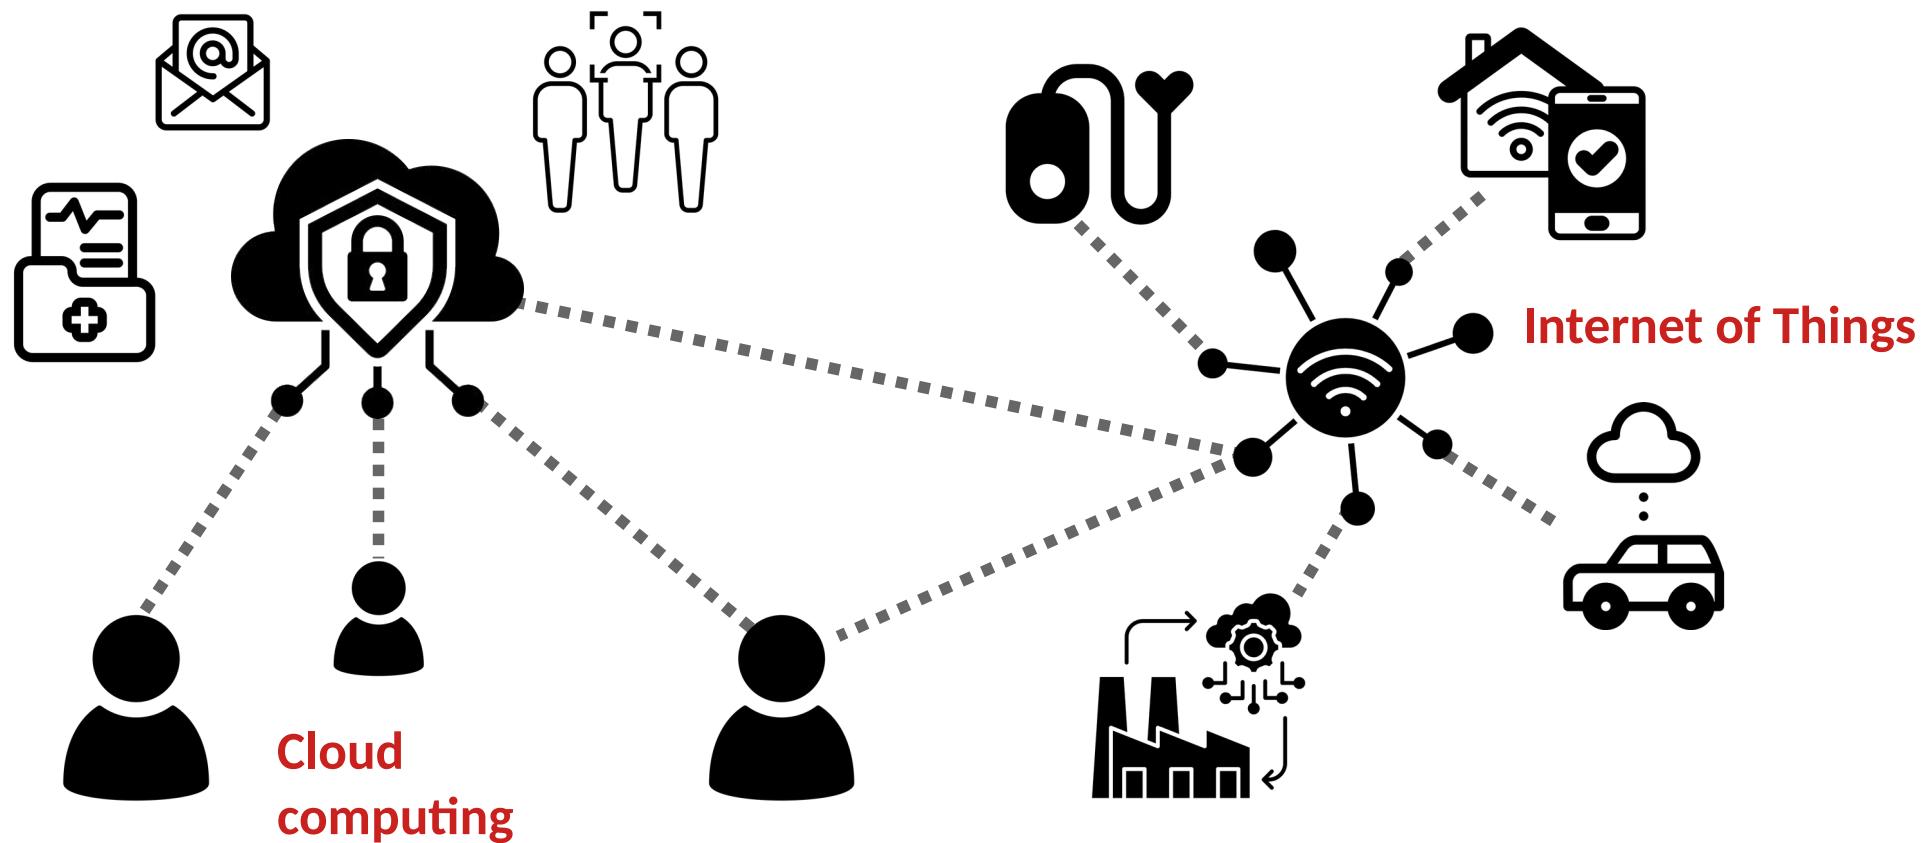


- 2004: ARM TrustZone
- 2015: **Intel Software Guard Extensions (SGX)**
- 2016: AMD Secure Encrypted Virtualization (SEV)
- 2018: IBM Protected Execution Facility (PEF)
- 2020: AMD SEV with Secure Nested Paging (SEV-SNP)
- 2022: Intel Trust Domain Extensions (TDX)
- 2023: ARM Confidential Compute Architecture (CCA)
- 2024: NVIDIA Confidential Computing



TEEs are here to stay...

“Confidential Computing Today, Just Computing Tomorrow” *



TEE Attack Research Leads the Way . . .



TEE Attack Research Leads the Way . . .



- Privileged TEE attacker models **sets the bar!**
- Idealized execution environment for attack research
- **Generalizations:** e.g., Foreshadow-NG, branch prediction, address translation, etc.

Motivation: Why Research TEE/SGX Security?



Overview About Intel News & Events Financial Info Stock Info Filings & Reports Board & Governance ESG

Overview

Press Releases

IR Calendar

Annual Stockholders' Meeting

Investor Meeting

Email Alerts

Presentations

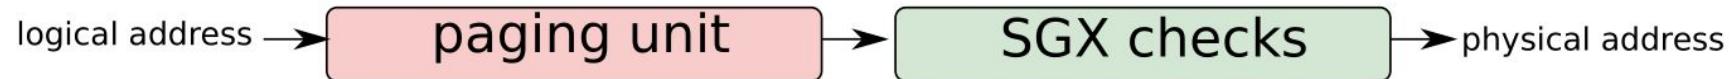
Data Protection across the Compute Stack

Technologies such as disk- and network-traffic encryption protect data in storage and during transmission, but data can be vulnerable to interception and tampering while in use in memory. "Confidential computing" is a rapidly emerging usage category that protects data while it is in use in a Trusted Execution Environment (TEE). Intel SGX is the most researched, updated and battle-tested TEE for data center confidential computing, with the smallest attack surface within the system. It enables application isolation in private memory regions, called enclaves, to help protect up to 1 terabyte of code and data while in use.



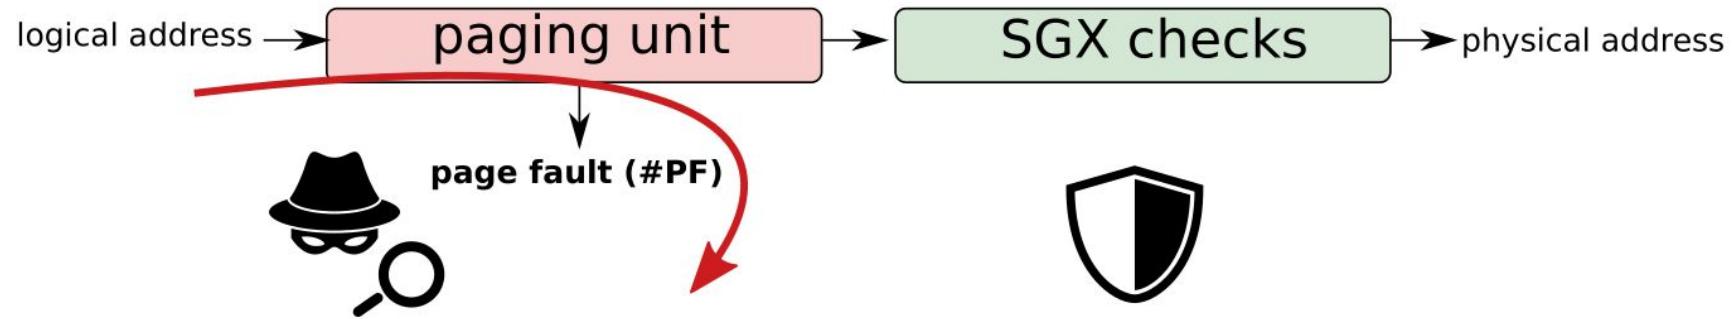
<https://www.intc.com/news-events/press-releases/detail/1423/intel-xeon-scalable-platform-built-for-most-sensitive>

Idea: Page Faults as a Side Channel



SGX machinery protects against direct address remapping attacks

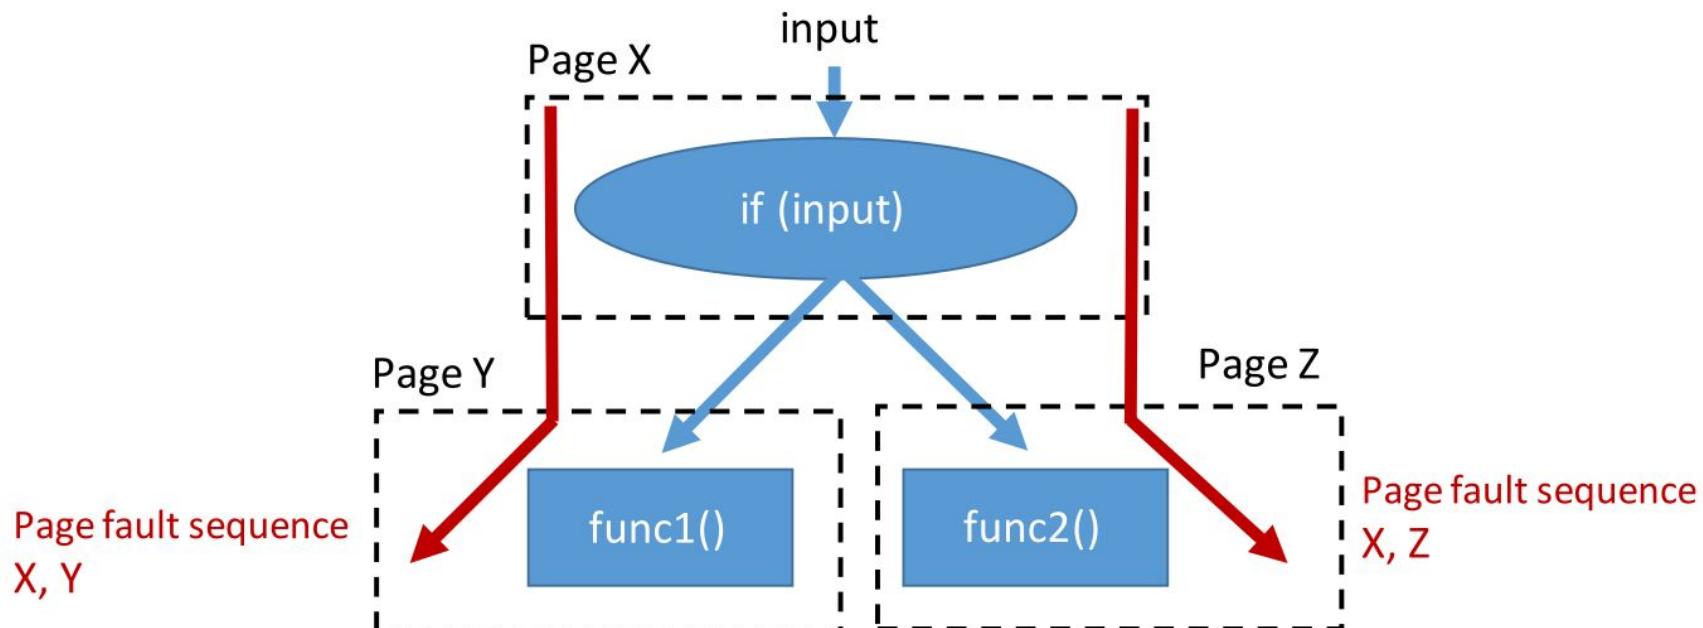
Idea: Page Faults as a Side Channel



... but untrusted address translation may **fault(!)**



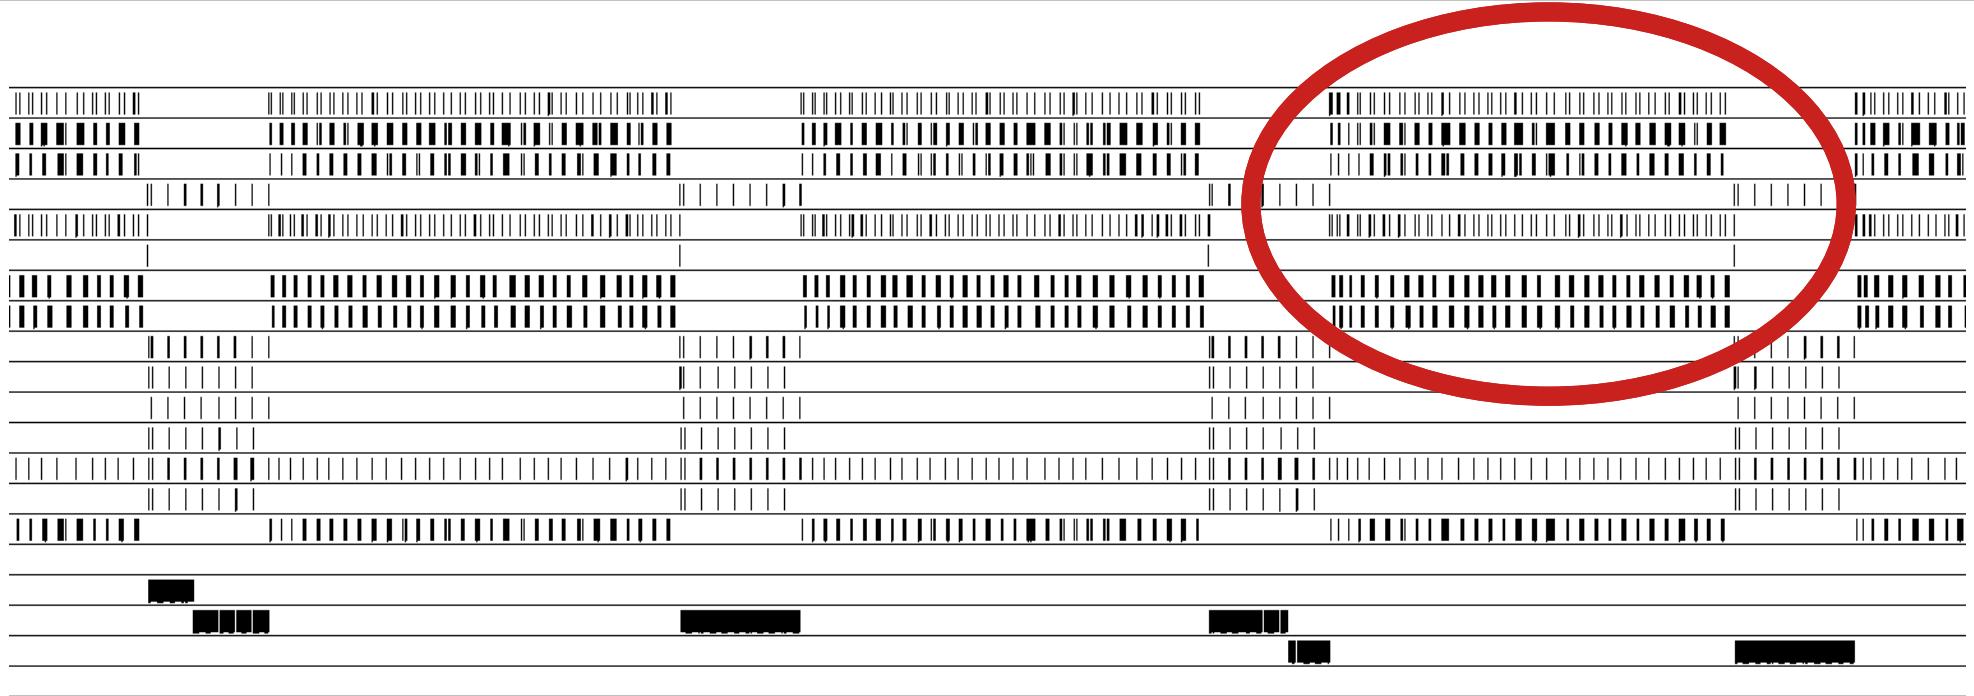
Intel SGX: Page Faults as a Side Channel



□ Xu et al.: “Controlled-channel attacks: Deterministic side channels for untrusted operating systems”, Oakland 2015.

⇒ Page fault traces leak **private control data/flow**

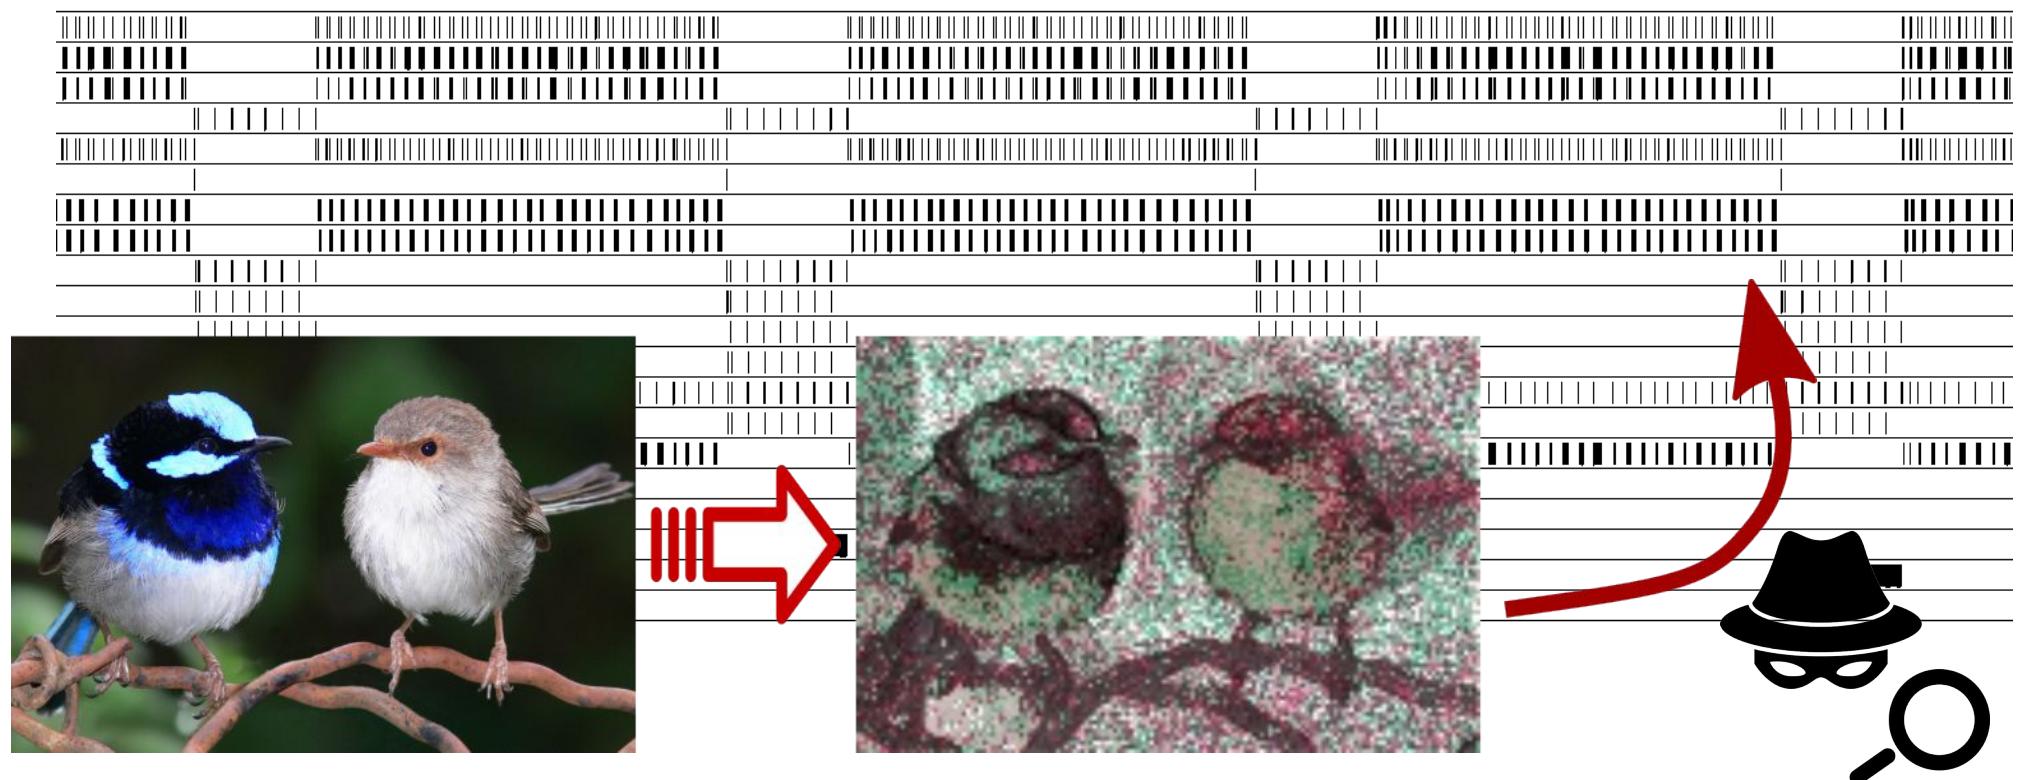
Spatial Resolution: Page-Granular Memory Access Traces

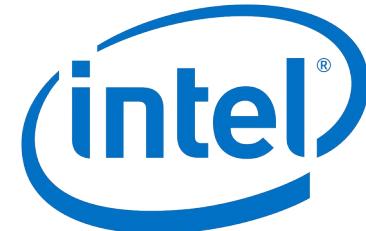


Detailed trace of (coarse-grained) code and data accesses over time...



Spatial Resolution: Page-Granular Memory Access Traces





Protection from Side-Channel Attacks

Intel® SGX does not provide explicit protection from side-channel attacks. It is the enclave developer's responsibility to address side-channel attack concerns.

In general, enclave operations that require an OCall, such as thread synchronization, I/O, etc., are exposed to the untrusted domain. If using an OCall would allow an attacker to gain insight into enclave secrets, then there would be a security concern. This scenario would be classified as a side-channel attack, and it would be up to the ISV to design the enclave in a way that prevents the leaking of side-channel information.

An attacker with access to the platform can see what pages are being executed or accessed. This side-channel vulnerability can be mitigated by aligning specific code and data blocks to exist entirely within a single page.

More important, the application enclave should use an appropriate crypto implementation that is side channel attack resistant inside the enclave if side-channel attacks are a concern.

Temporal Resolution Limitations for the Page-Fault Oracle

```
1 size_t strlen (char *str)
2 {
3     char *s;
4
5     for (s = str; *s; ++s);
6     return (s - str);
7 }
```

```
1      mov   %rdi,%rax
2 1:  cmpb $0x0 ,(%rax)
3     je    2f
4     inc   %rax
5     jmp   1b
6 2:  sub   %rdi,%rax
7     retq
```

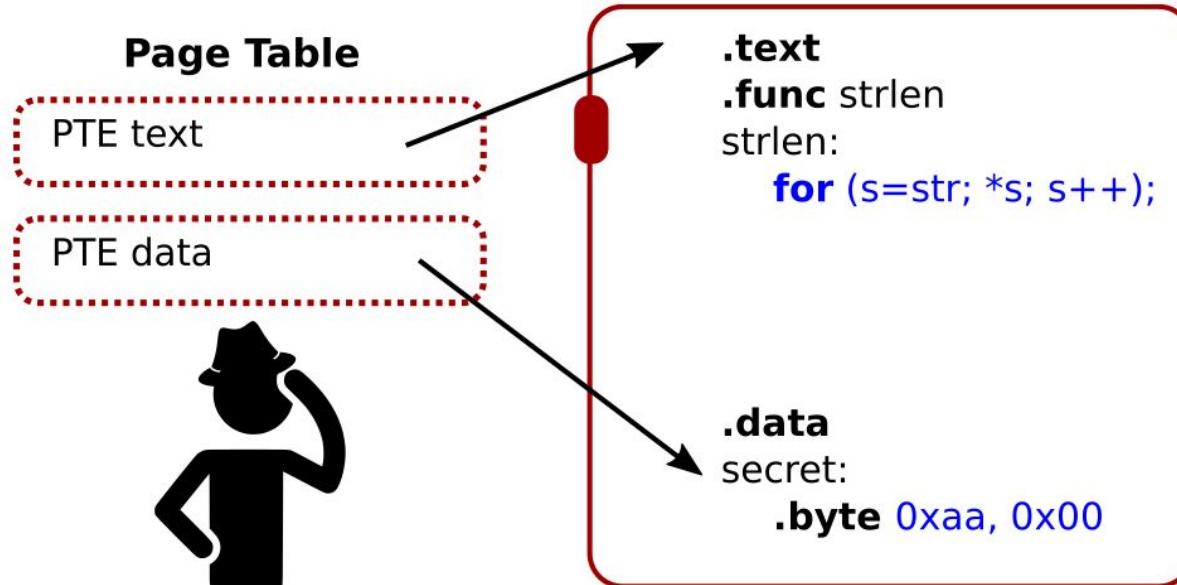
⇒ tight loop: 4 instructions, single memory operand, single code + data page

Counting strlen loop iterations?



Note: Page-fault attacks cannot make progress for 1 code + data page

Temporal Resolution Limitations for the Page-Fault Oracle



Counting `strlen` loop iterations?

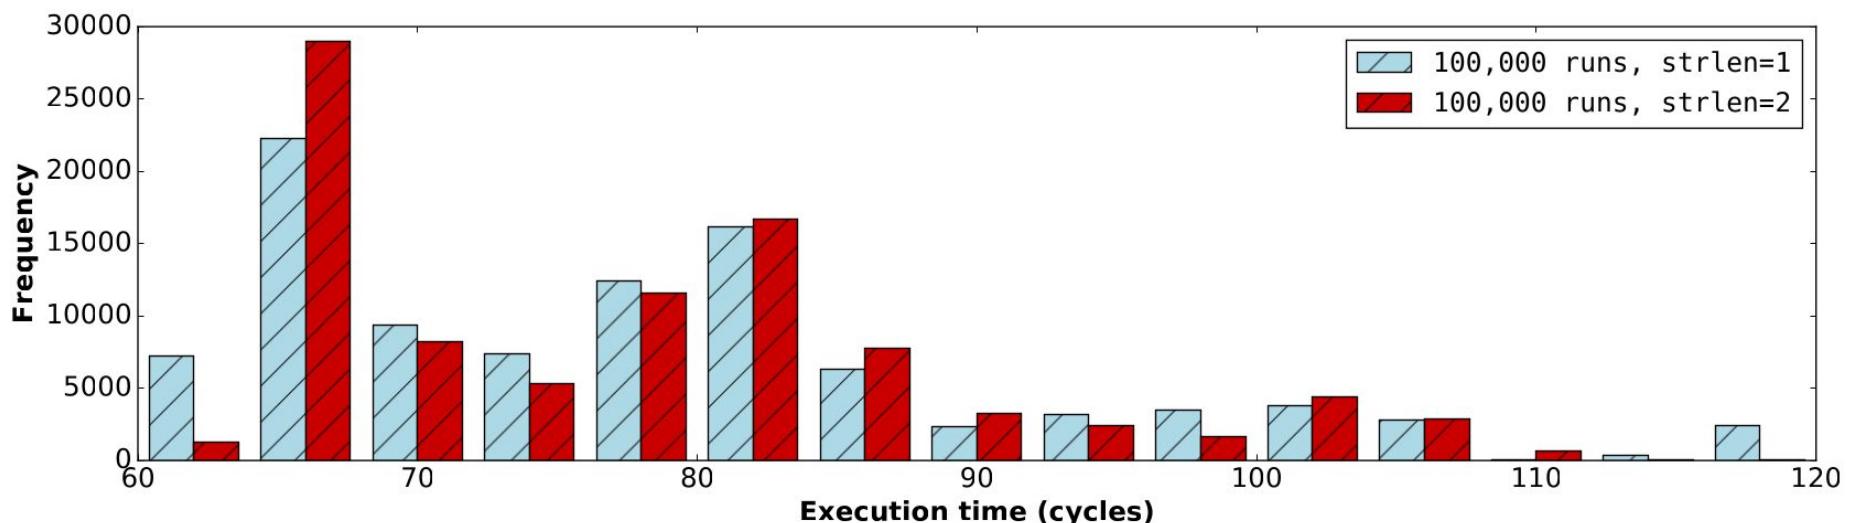


Progress requires both pages present (non-faulting) \leftrightarrow page fault oracle

Building the Side-Channel Oracle with Execution Timing?



Too noisy: modern x86 processors are lightning fast...



Challenge: Side-Channel Sampling Rate



Slow
shutter speed

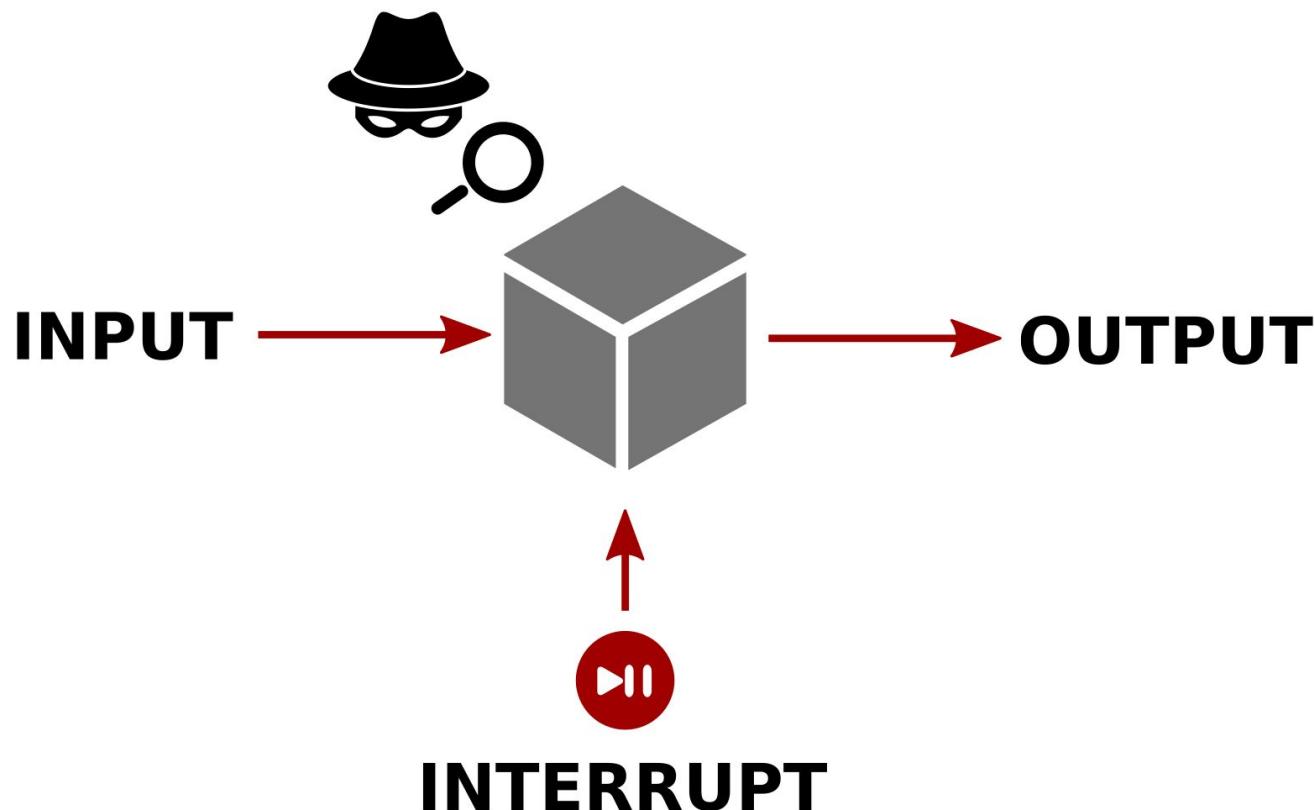


Medium
shutter speed



Fast
shutter speed

SGX-Step: Executing Enclaves one Instruction at a Time



SGX-Step: Executing Enclaves one Instruction at a Time



SGX-Step



<https://github.com/jovanbulck/sgx-step>

Watch

22

Star

245

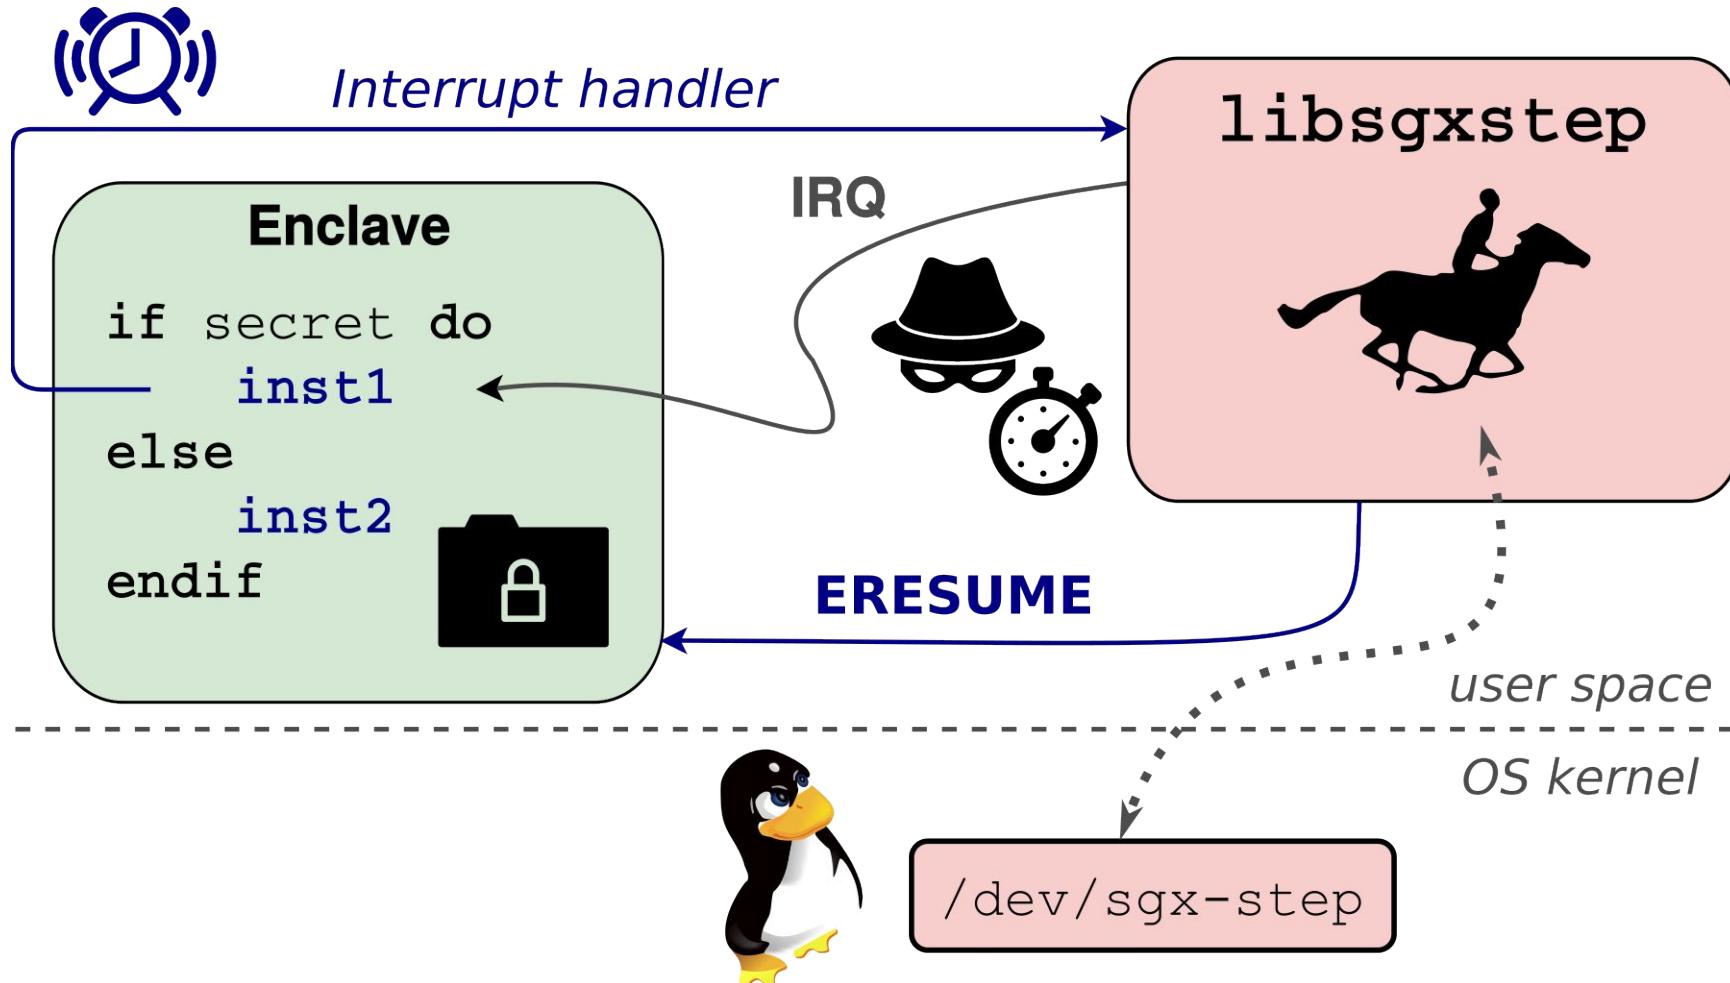
Fork

52



Van Bulck et al., "SGX-Step: A Practical Attack Framework for Precise Enclave Execution Control", SysTEX 2017.

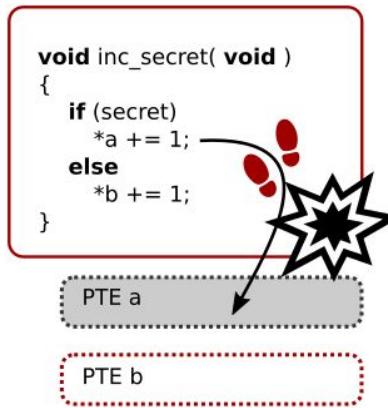
SGX-Step: Executing Enclaves one Instruction at a Time



SGX-Step Demo: Single-Stepping Password Comparison

```
jo@breuer:~/sgx-step-demo$ sudo ./app
```

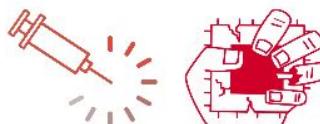
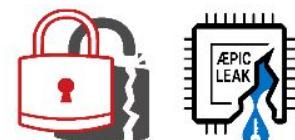

SGX-Step: A Versatile Open-Source Attack Framework



Page-table manipulation



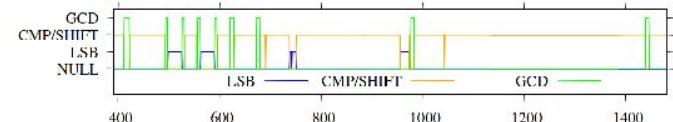
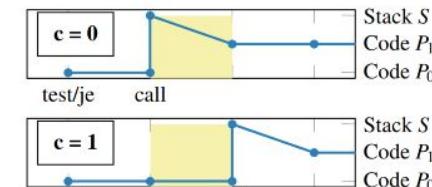
Interrupt latency



High-resolution probing



SGX-Step

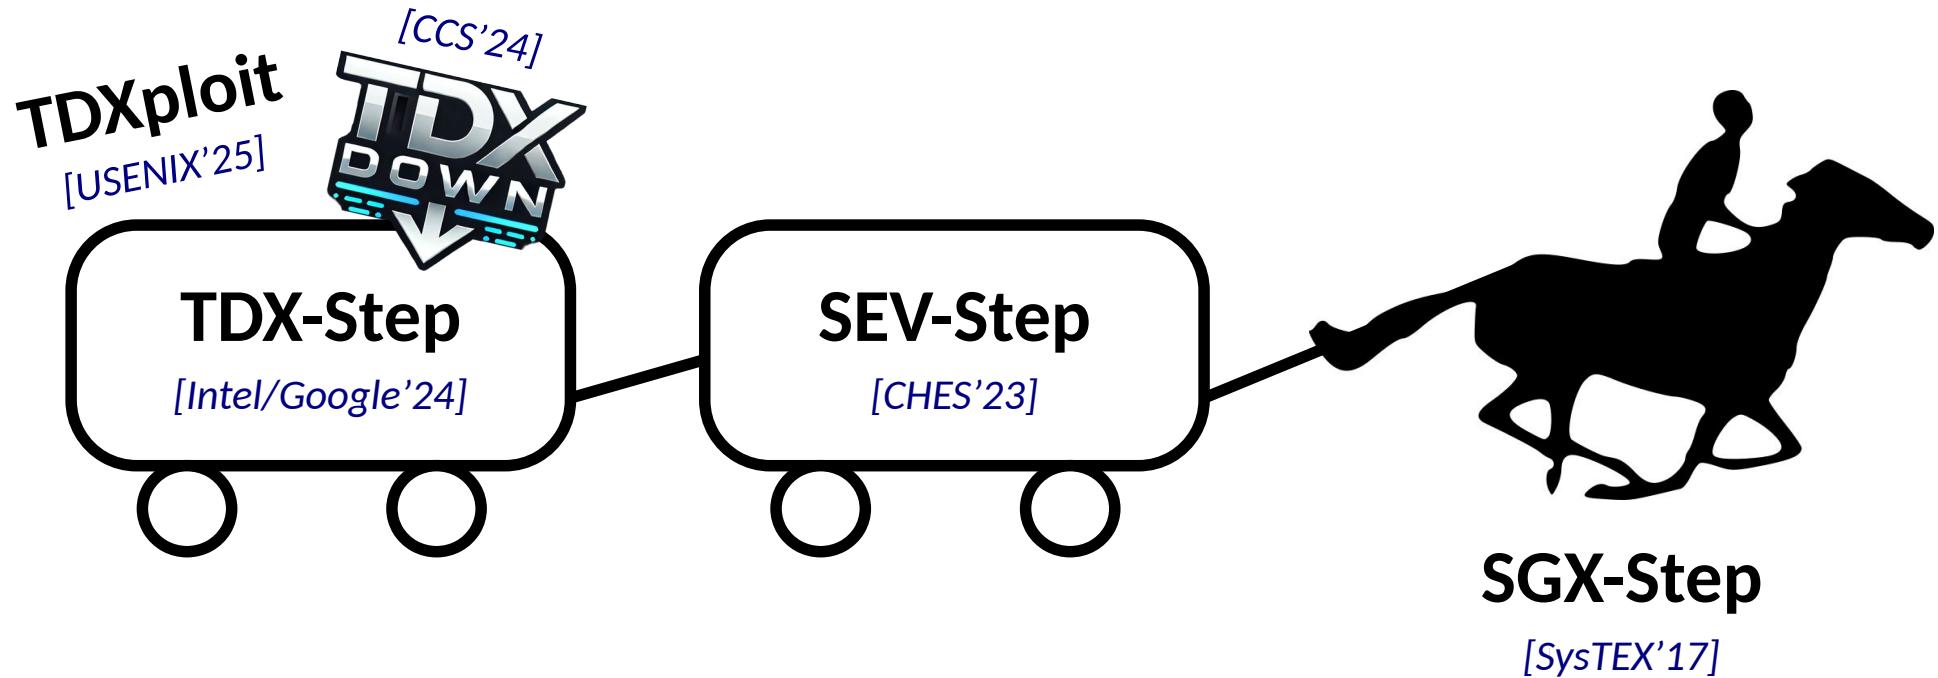


Interrupt counting

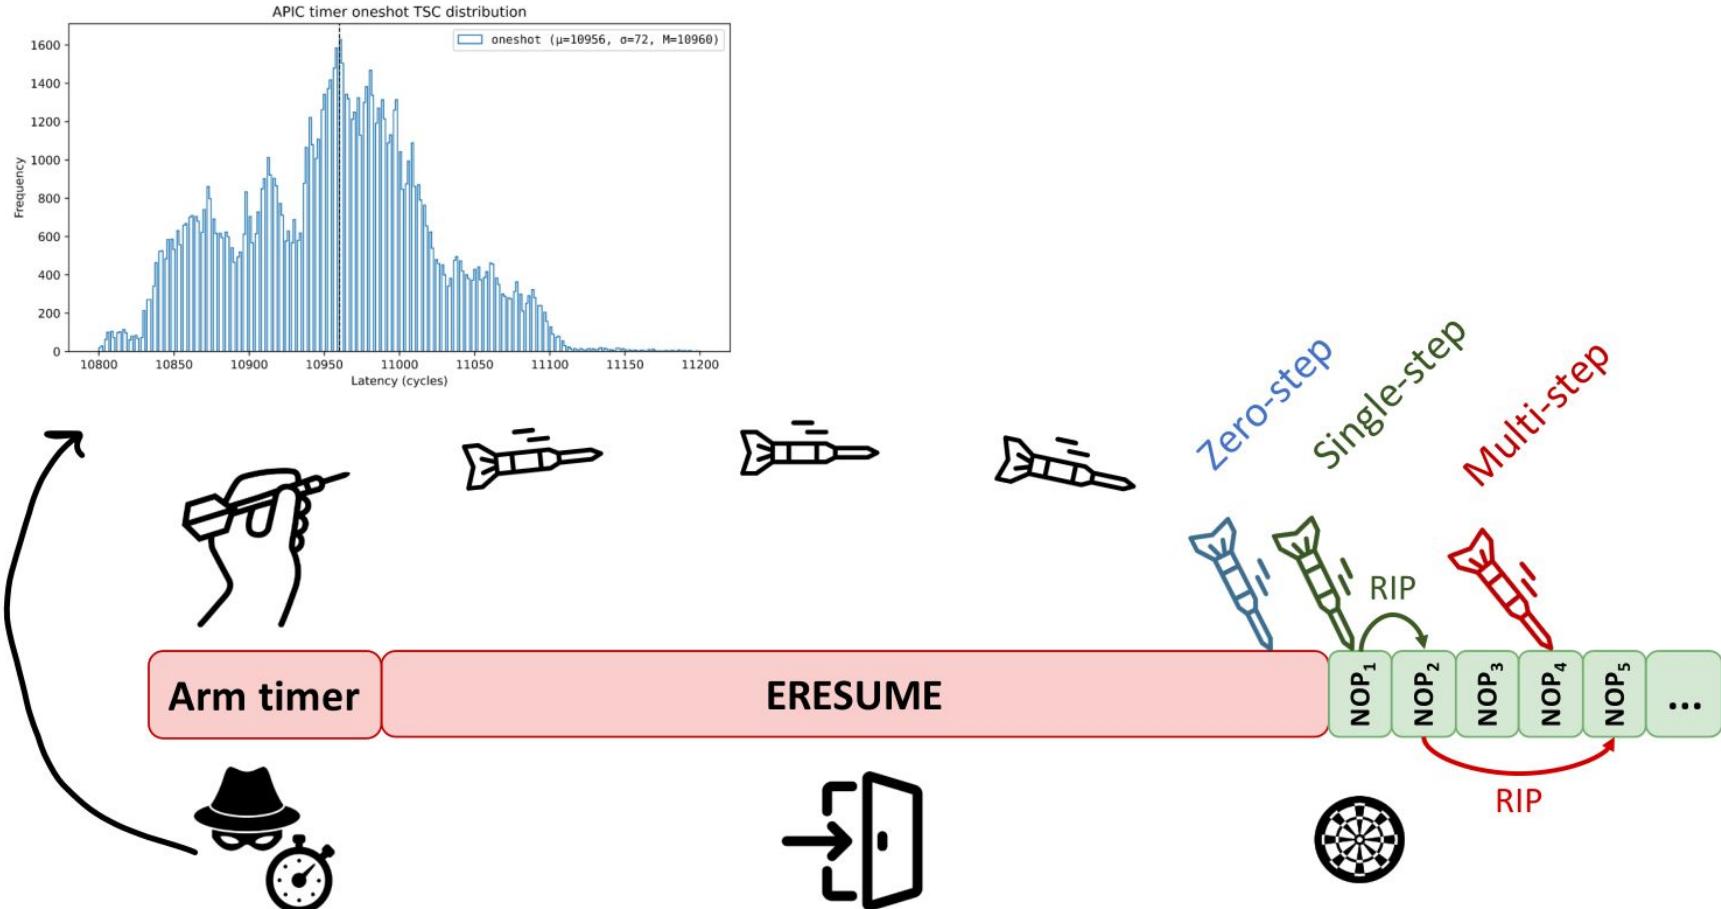


Zero-step replaying

Single-Stepping Beyond Intel SGX



Root-causing SGX-Step: Aiming the timer interrupt



Root-causing SGX-Step: Microcode assists to the rescue!

PTE A-bit	Mean (cycles)	Stddev (cycles)
A=1	27	30
A=0	666	55



3. Assisted PT walk



1. Clear PTE A-bit



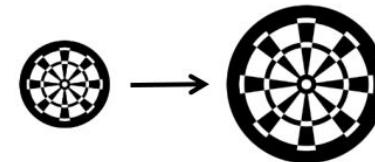
2. TLB flush



Arm timer

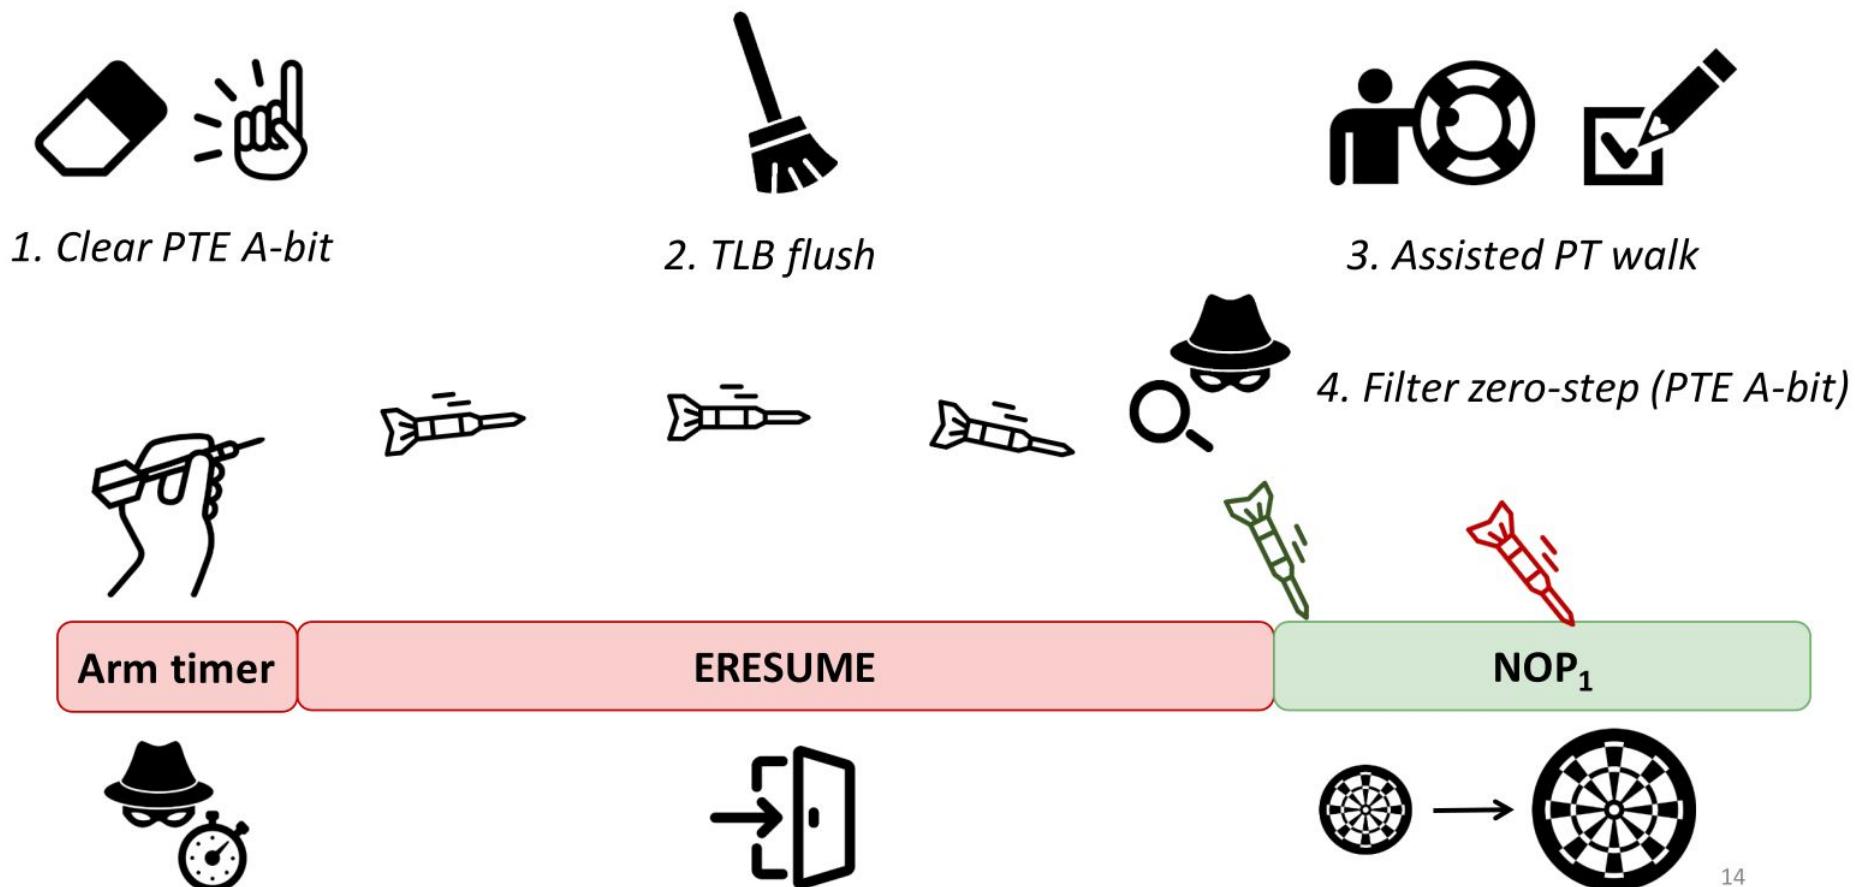
ERESUME

NOP₁

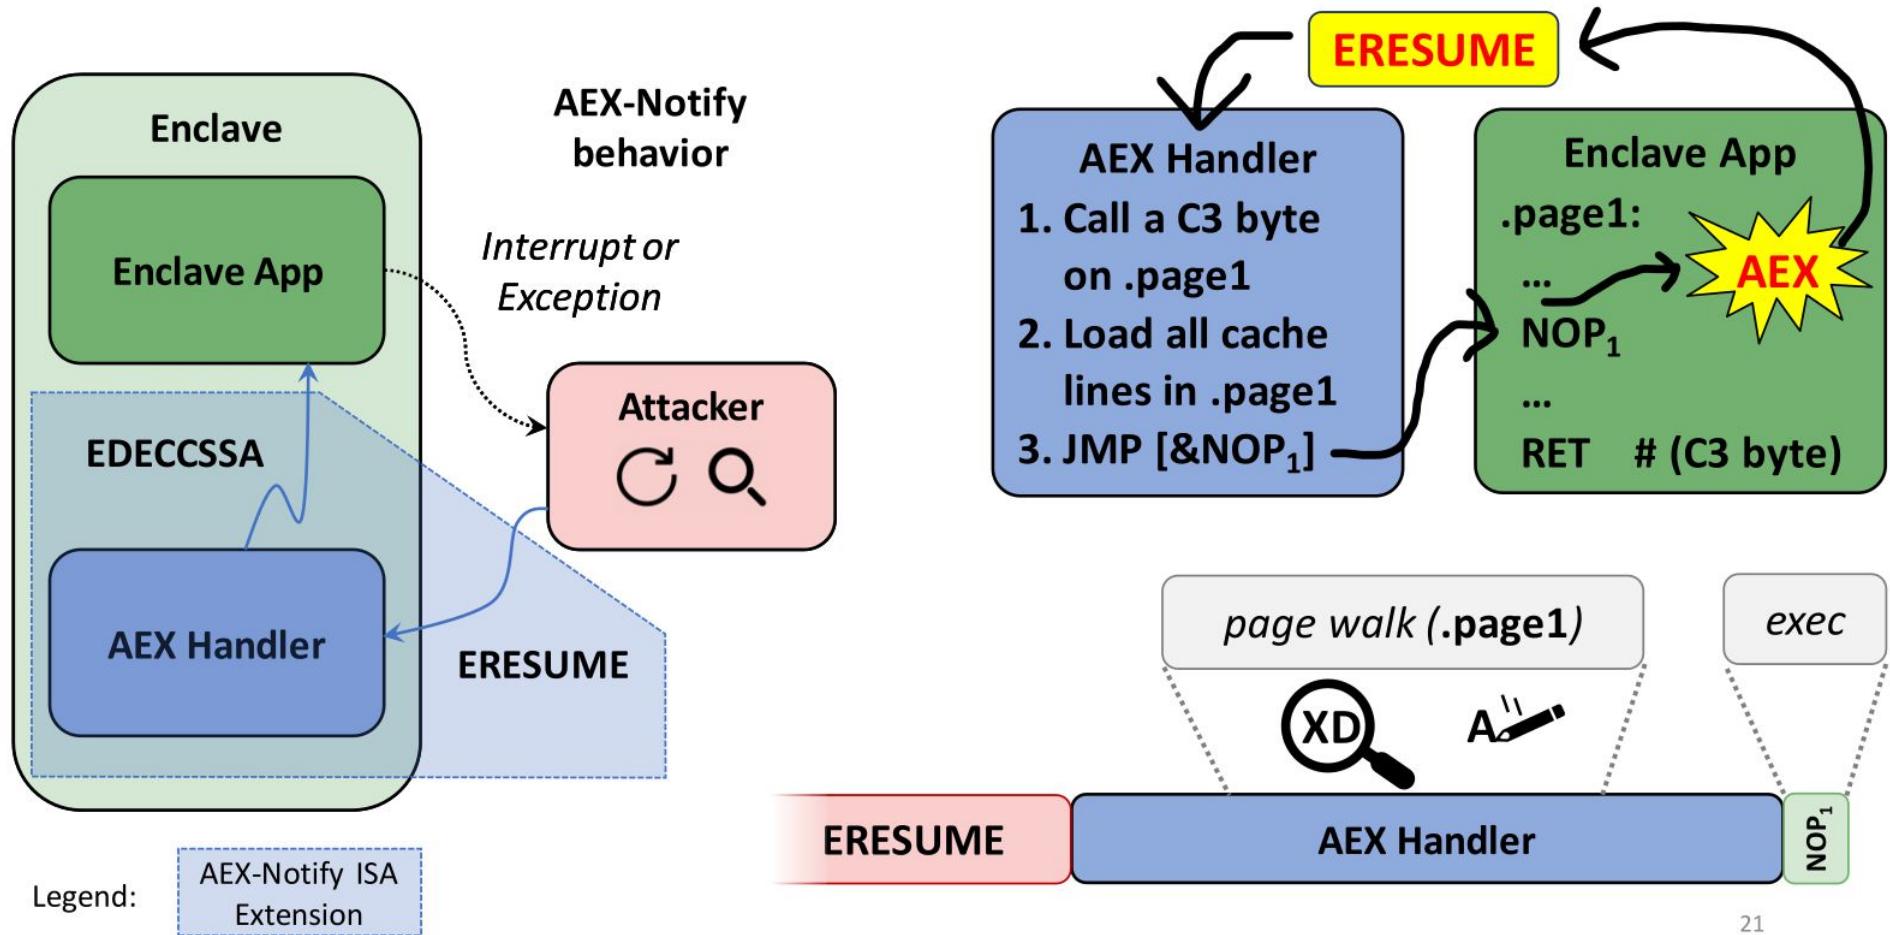


12

Root-causing SGX-Step: Microcode assists to the rescue!

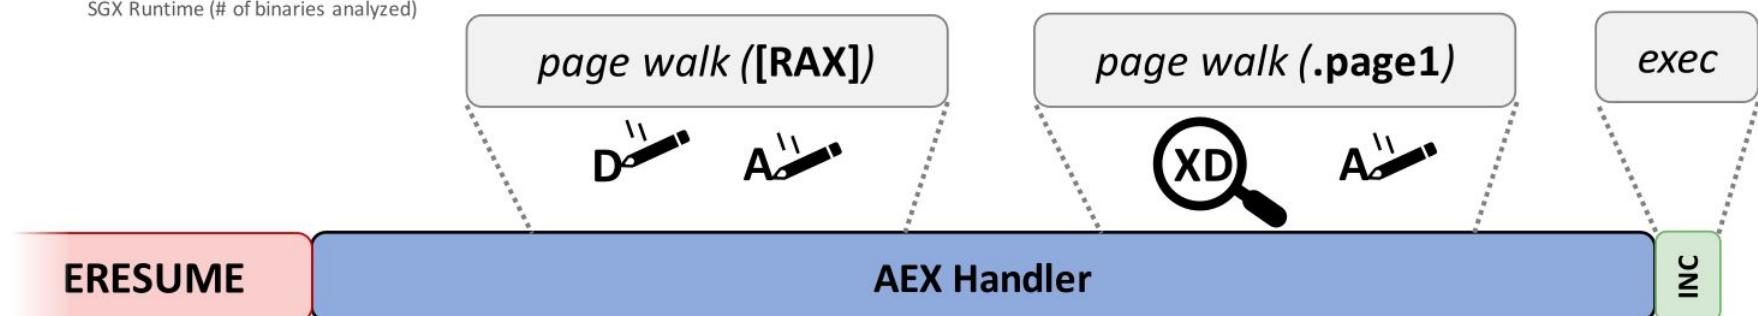
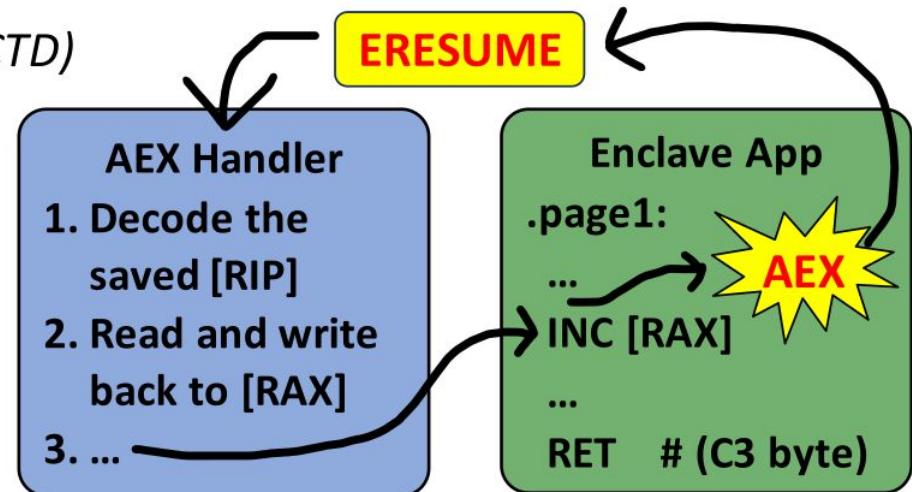
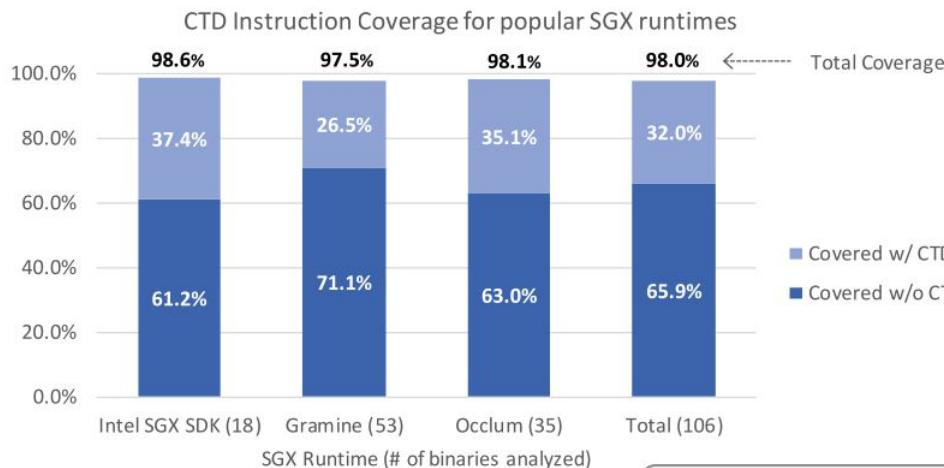


AEX-Notify: Hardware-Software Co-Design Solution



AEX-Notify: Hardware-Software Co-Design Solution

We implemented a fast, constant-time decoder (CTD)



CHAPTER 8

ASYNCHRONOUS ENCLAVE EXIT NOTIFY AND THE EDECCSSA USER LEAF FUNCTION

8.1 INTRODUCTION

Asynchronous Enclave Exit Notify (AEX-Notify) is an extension to Intel® SGX that allows Intel SGX enclaves to be notified after an asynchronous enclave exit (AEX) has occurred. EDECCSSA is a new Intel SGX user leaf function (ENCLU[EDECCSSA]) that can facilitate AEX notification handling, as well as software exception handling. This chapter provides information about changes to the Intel SGX architecture that support AEX-Notify and ENCLU[EDECCSSA].

The following list summarizes the a details are provided in Section 8.3)

- SECS.ATTRIBUTES.AEXNOTIFY
- TCS.FLAGS.AEXNOTIFY: This en
- SSA.GPRSGX.AEXNOTIFY: Enclave-writable byte that allows enclave software to dynamically enable/disable AEX notifications.



SGX-Step led to new x86 processor instructions!

→ shipped in millions of devices \geq 4th Gen Xeon CPU

An AEX notification is delivered by ENCLU[ERESUME] when the following conditions are met:

Intel AEX Notify Support Prepped For Linux To Help Enhance SGX Enclave Security

Written by Michael Larabel in Intel on 6 November 2022 at 06:01 AM EST. [5 Comments](#)



Future Intel CPUs and some existing processors via a microcode update will support a new feature called the Asynchronous EXit (AEX) notification mechanism to help with Software Guard Extensions (SGX) enclave security. Patches for the Linux kernel are pending for implementing this Intel AEX Notify support with capable processors.

Intel's Asynchronous EXit (AEX) notification mechanism lets SGX enclaves run a handler after an AEX event. Those handlers can be used for things like mitigating SGX-Step as an attack framework for precise enclave execution control.



Code 1 in intel/linux-sgx X

intel sdk/trts/linux/trts_mitigation.S

```
48 * Description:  
49 *      The file provides mitigations for SGX-Step  
50 */  
51 * Function:  
52 *      constant_time_apply_sgxstep_mitigation_and_continue_execution  
53 *      Mitigate SGX-Step and return to the point at which the  
most recent  
54 *      interrupt/exception occurred.
```

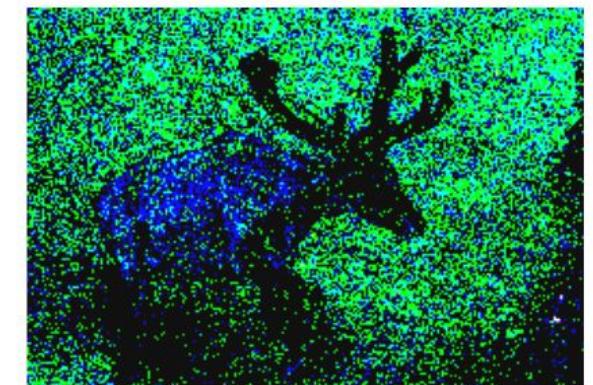
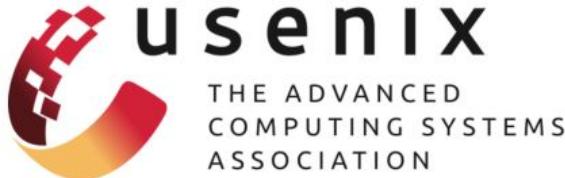


SGX-Step led to changes in major OSs and enclave SDKs

There's a Catch...

Finally note that our proposed mitigation does not protect against interrupting enclaves and observing application code and data page accesses at a coarse-grained 4 KiB spatial resolution. In contrast to the fine-grained, instruction-granular interrupt-driven attacks we consider in this work, such controlled-channel attacks have received ample attention [18, 47, 56, 59] from the research community.

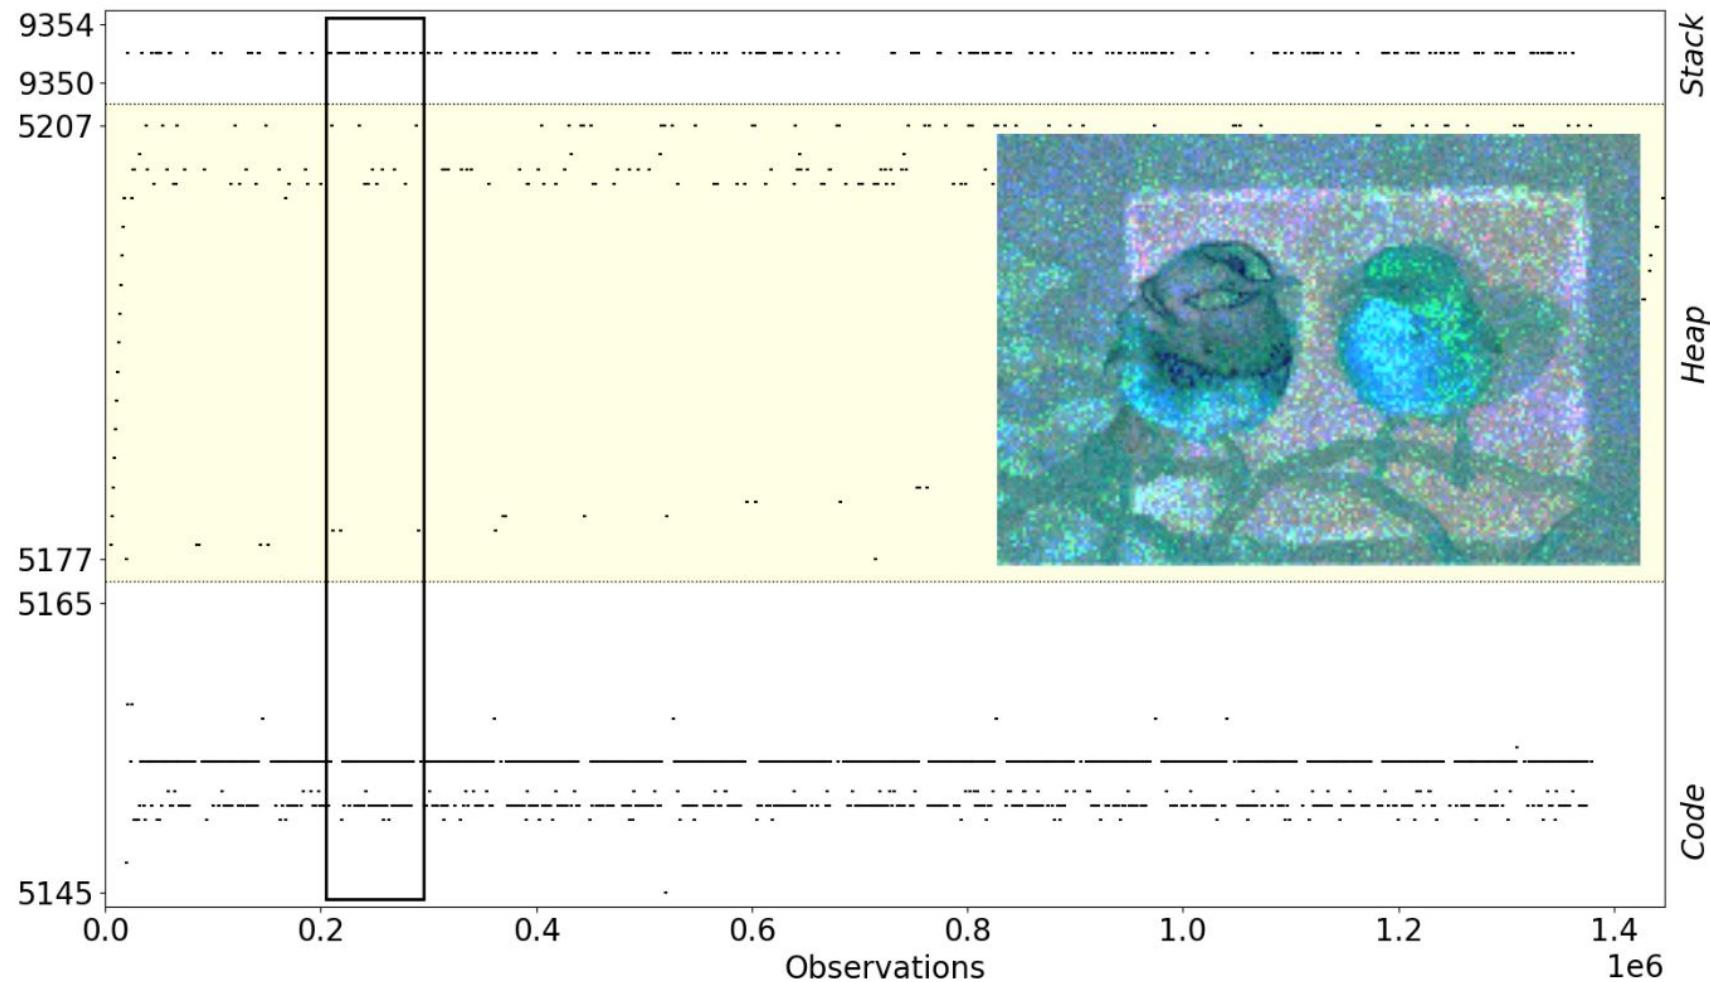
Why Mitigating Single-Stepping is Not Enough



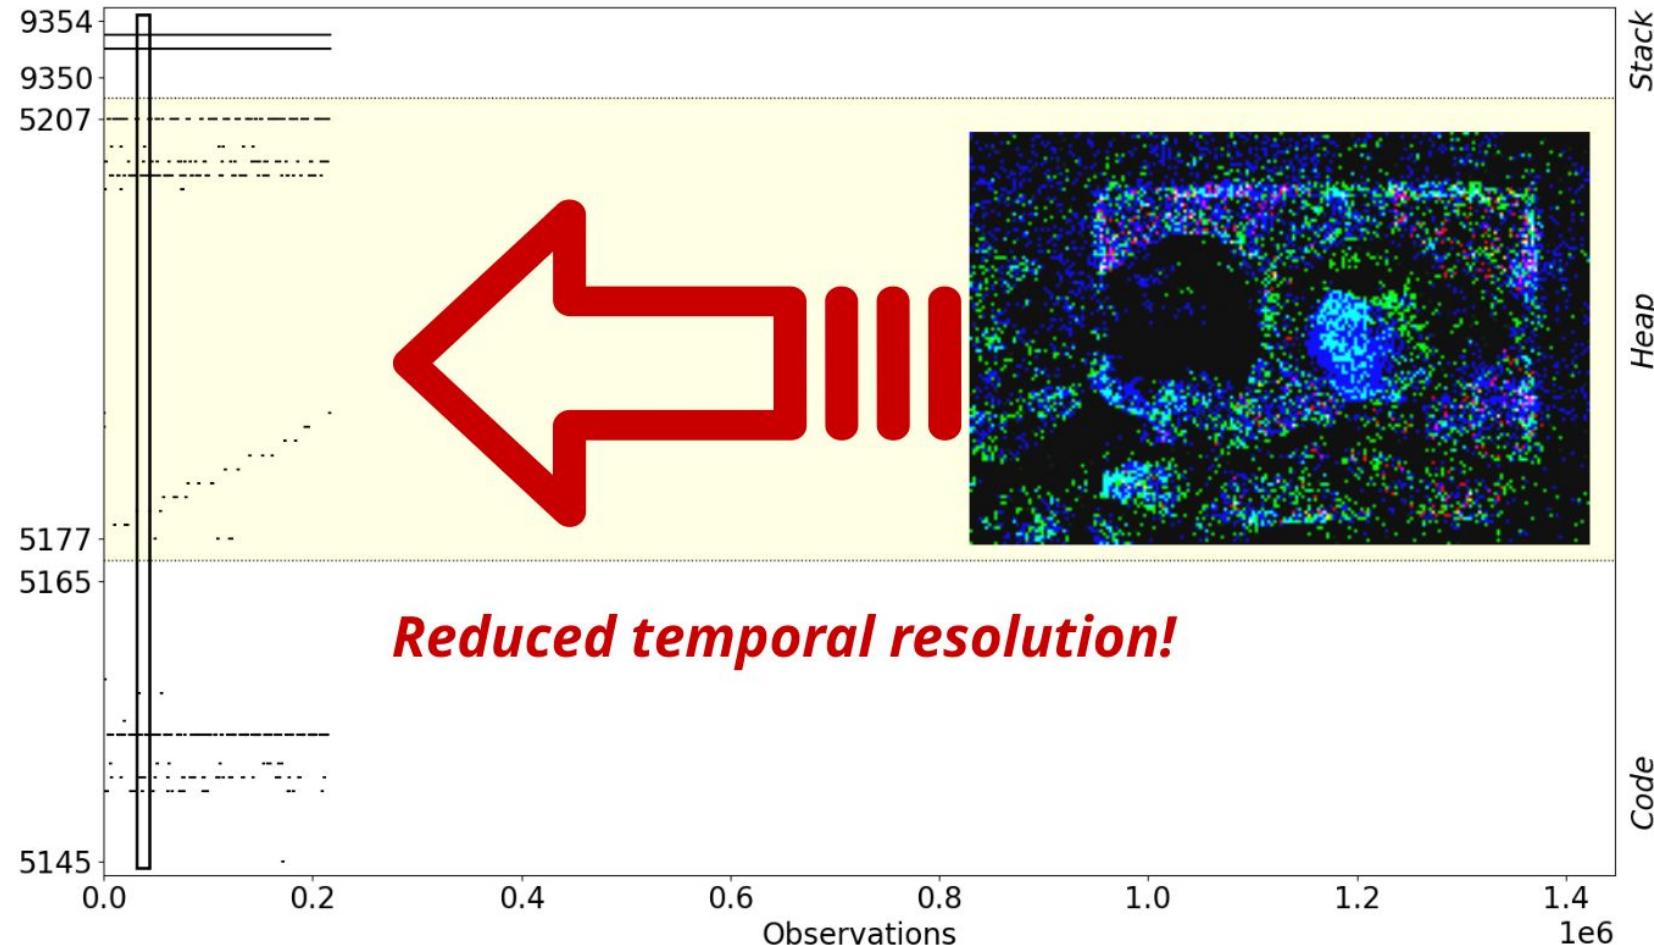
Original (left), Xu et al. (middle), our attack with AEX-Notify single-stepping defense (right)



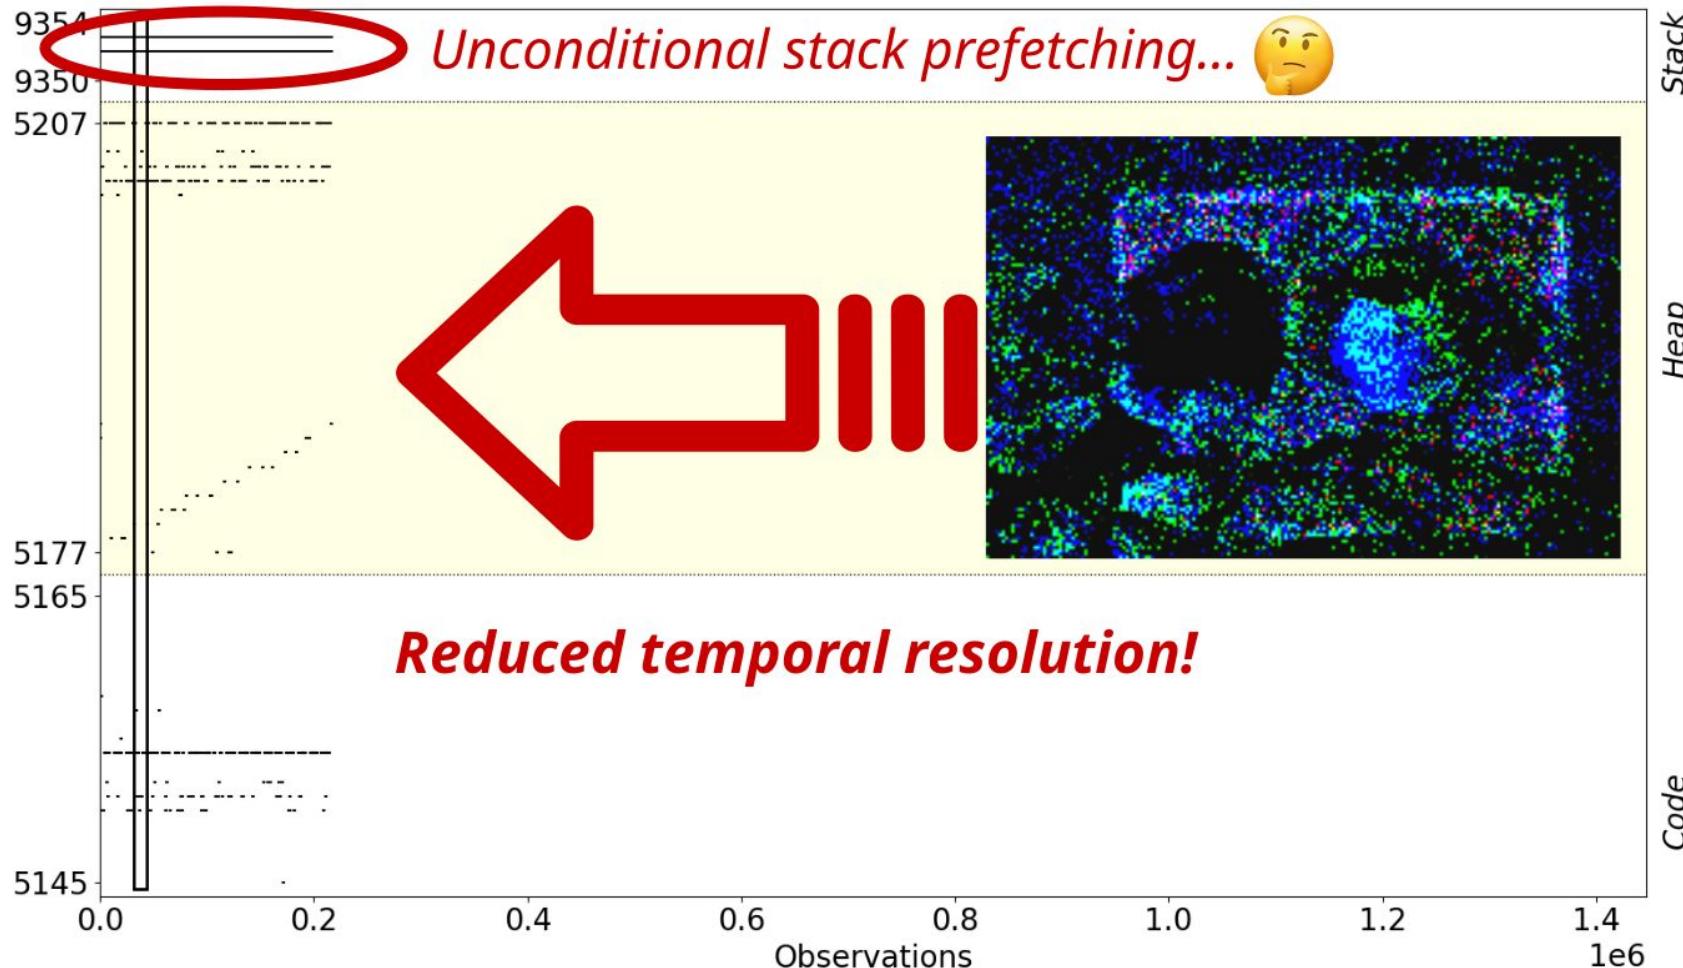
Libjpeg: AEX-Notify's Temporal Reduction in Practice



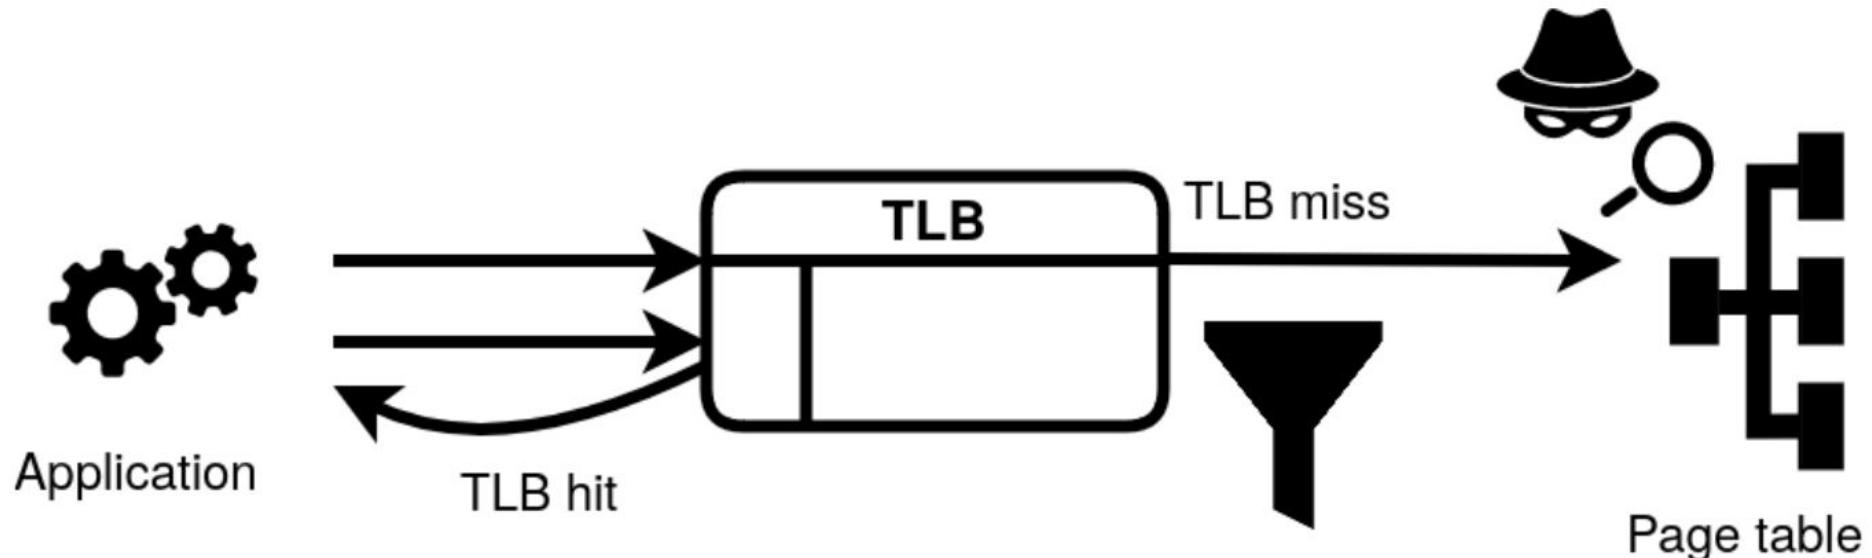
Libjpeg: AEX-Notify's Temporal Reduction in Practice



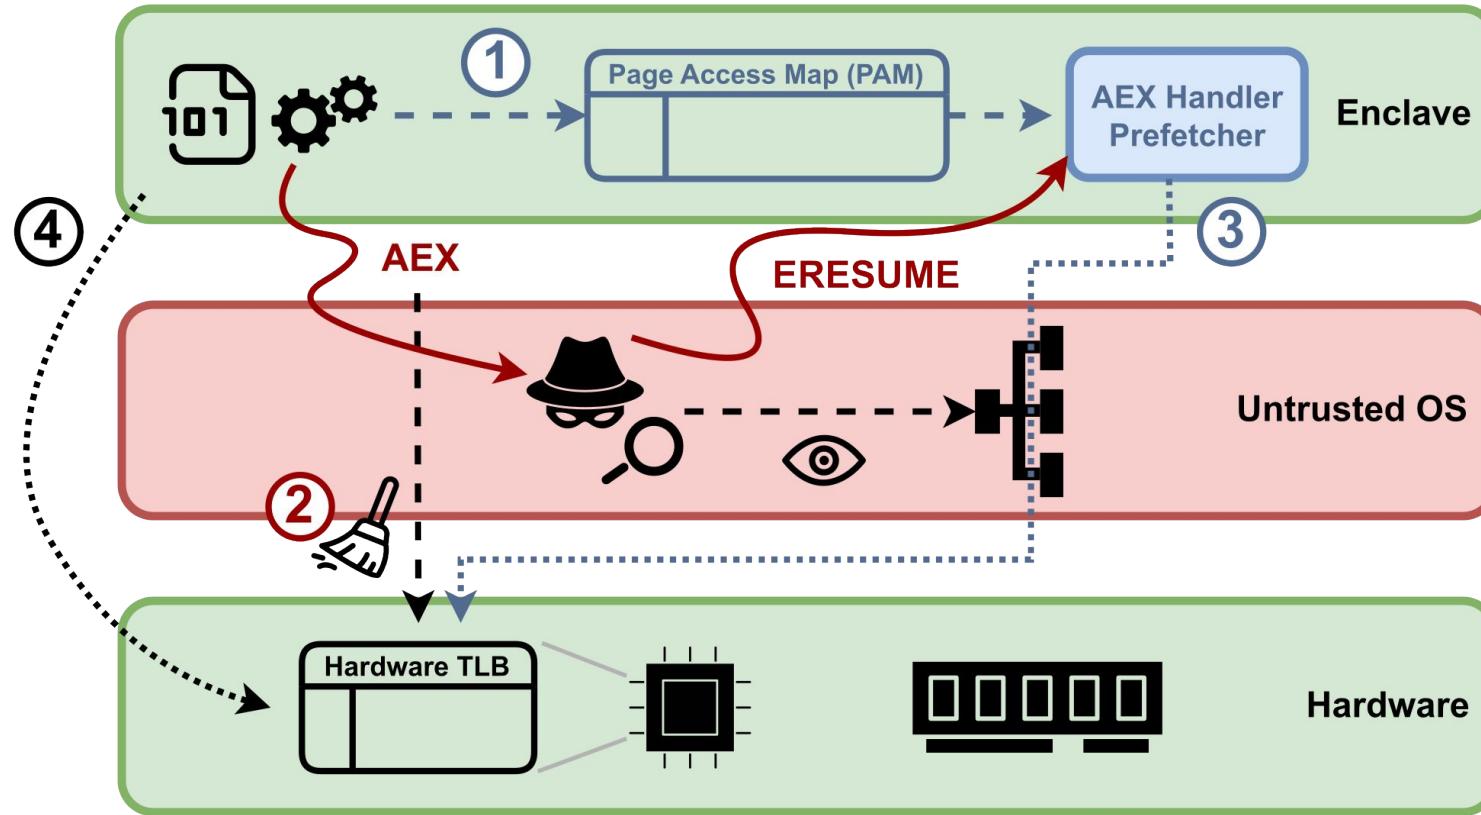
Libjpeg: AEX-Notify's Temporal Reduction in Practice



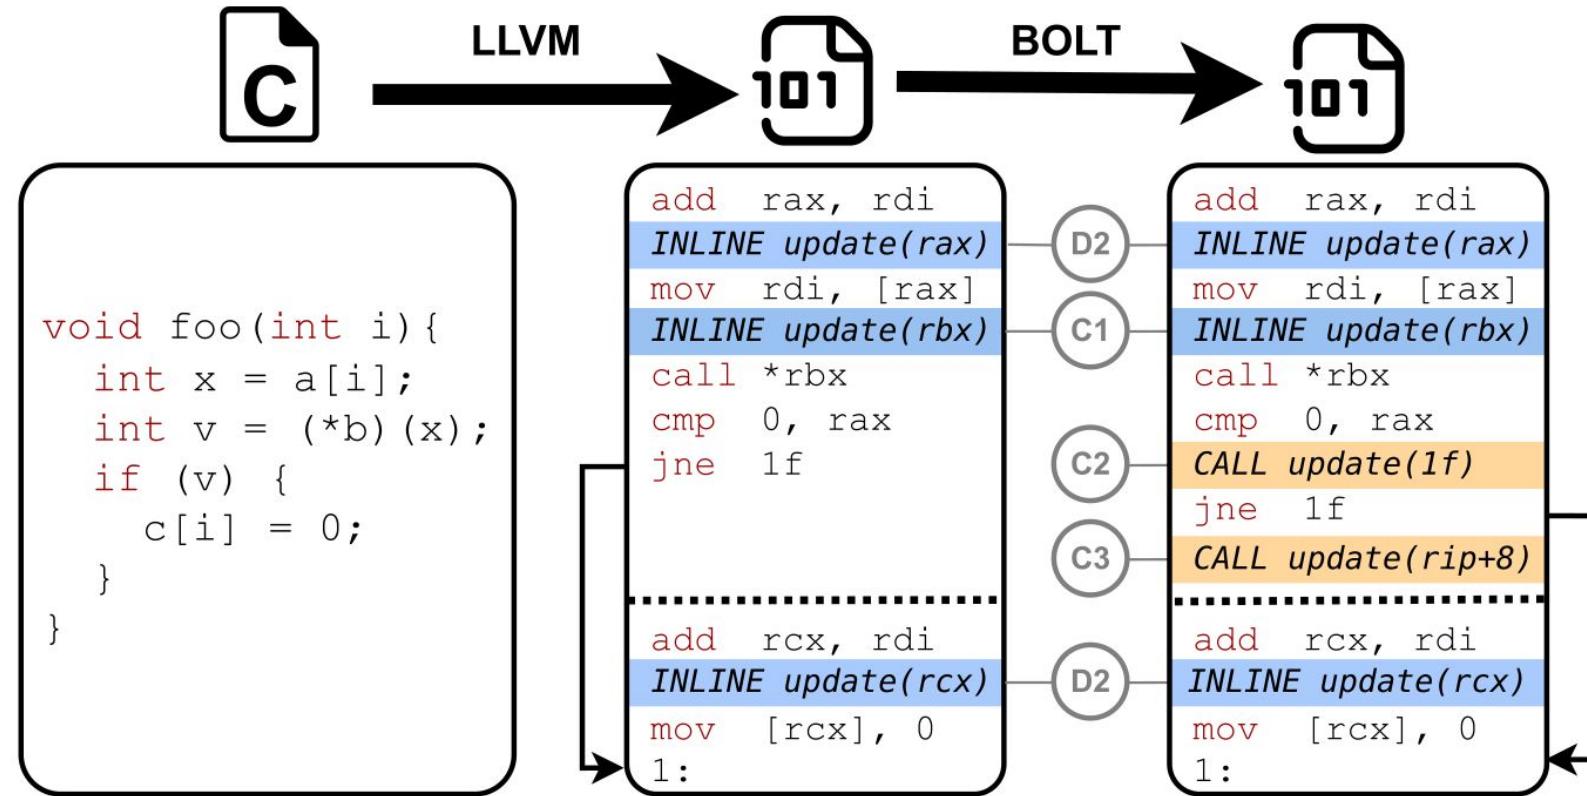
Idea: TLB as a “Filter” to Hide Page Accesses



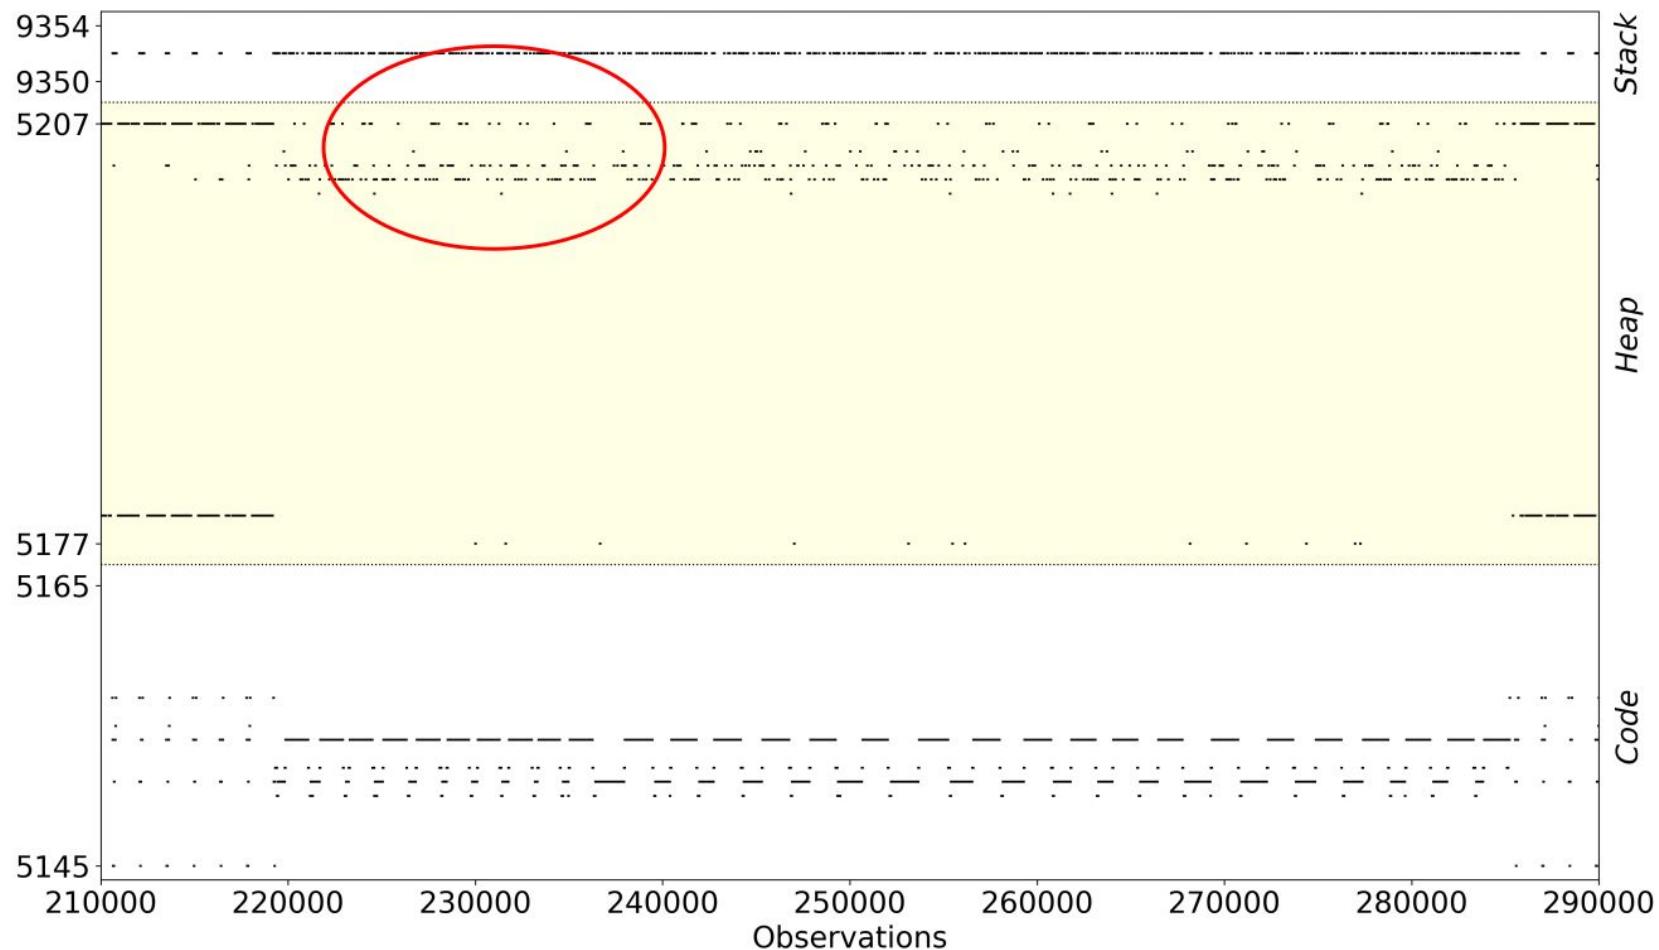
TLBlur: Self-Monitoring and Restoring Enclave Page Accesses



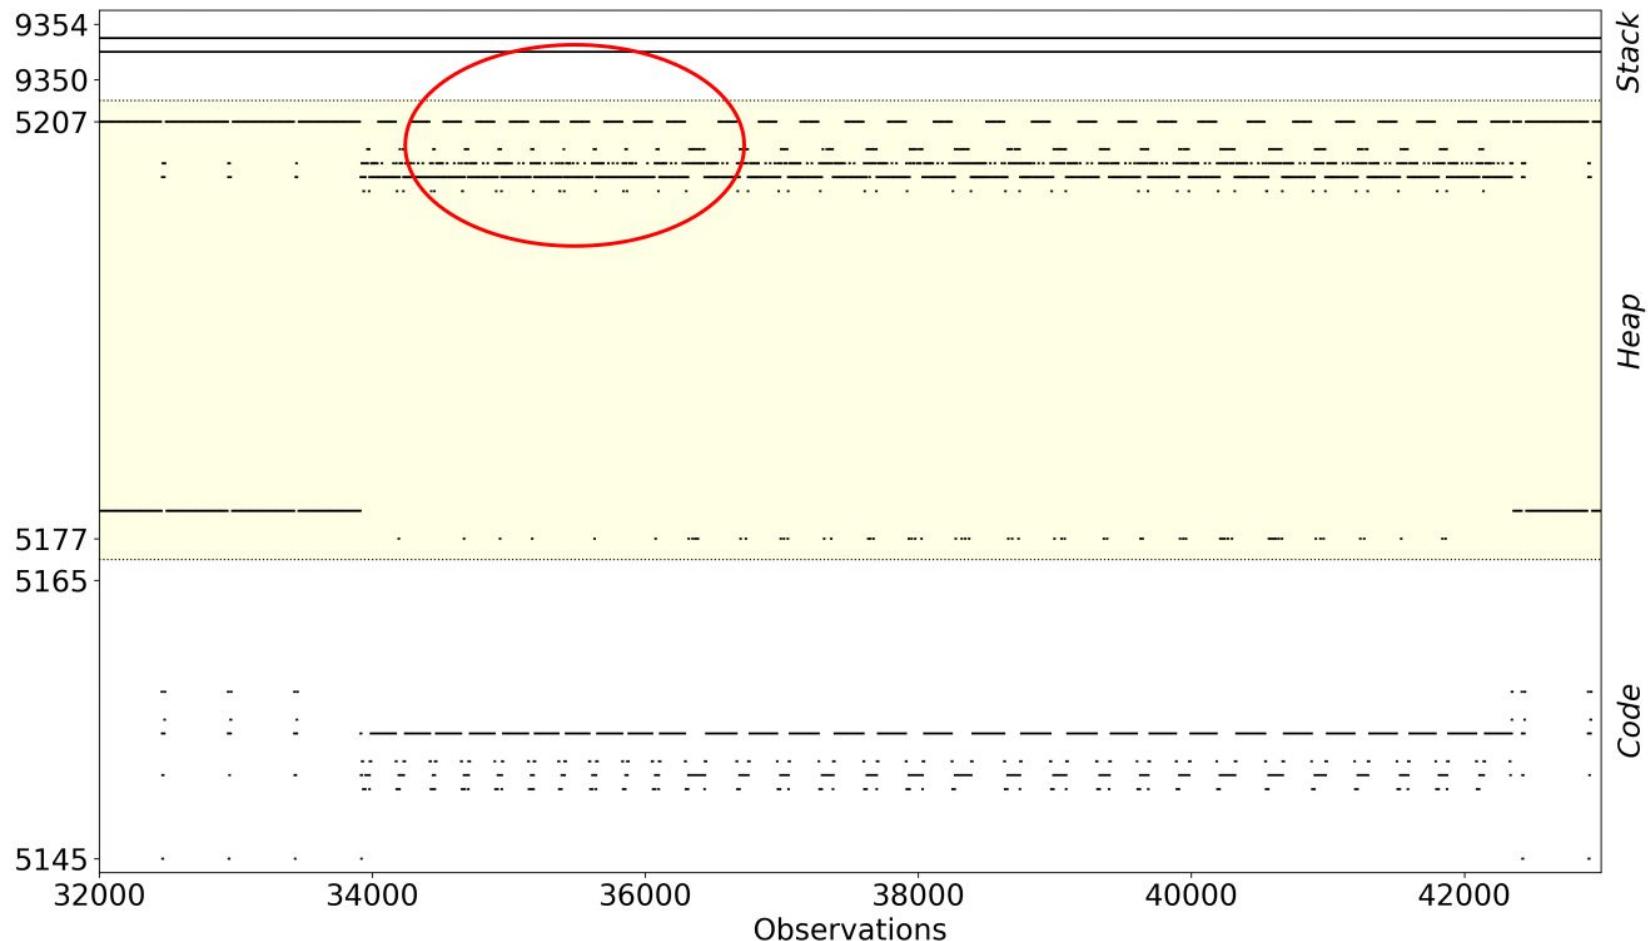
Instrumentation to Self-Monitor Page Accesses at Runtime



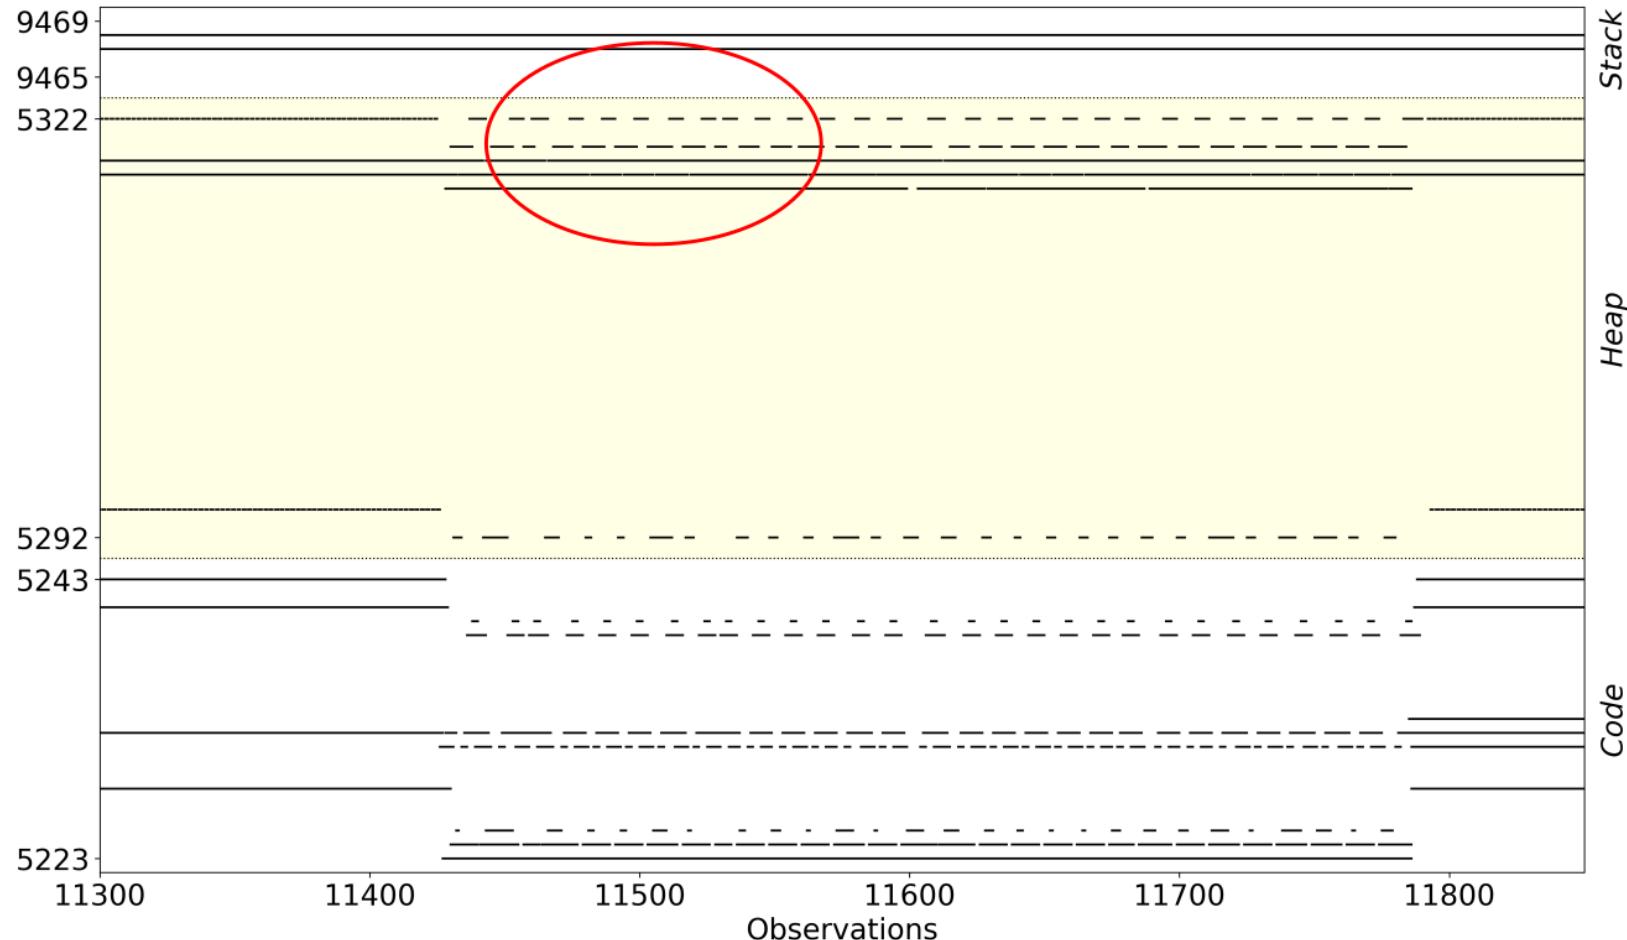
Leakage Reduction in Practice: Libjpeg Single-Stepping



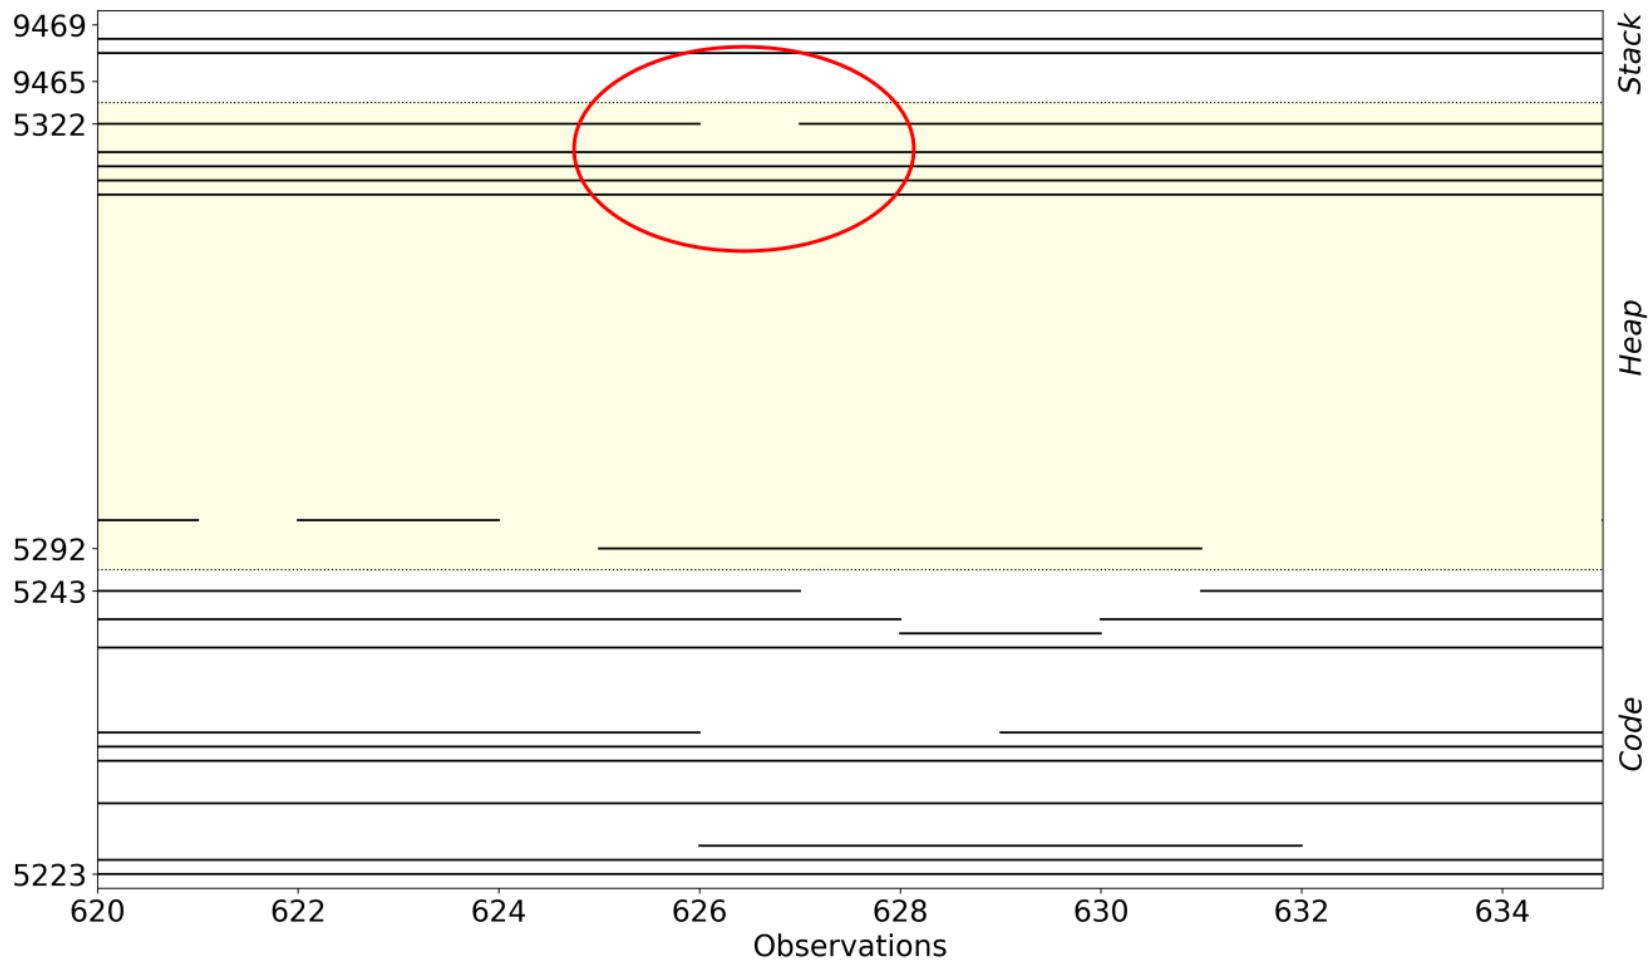
Leakage Reduction in Practice: Libjpeg Page Faults



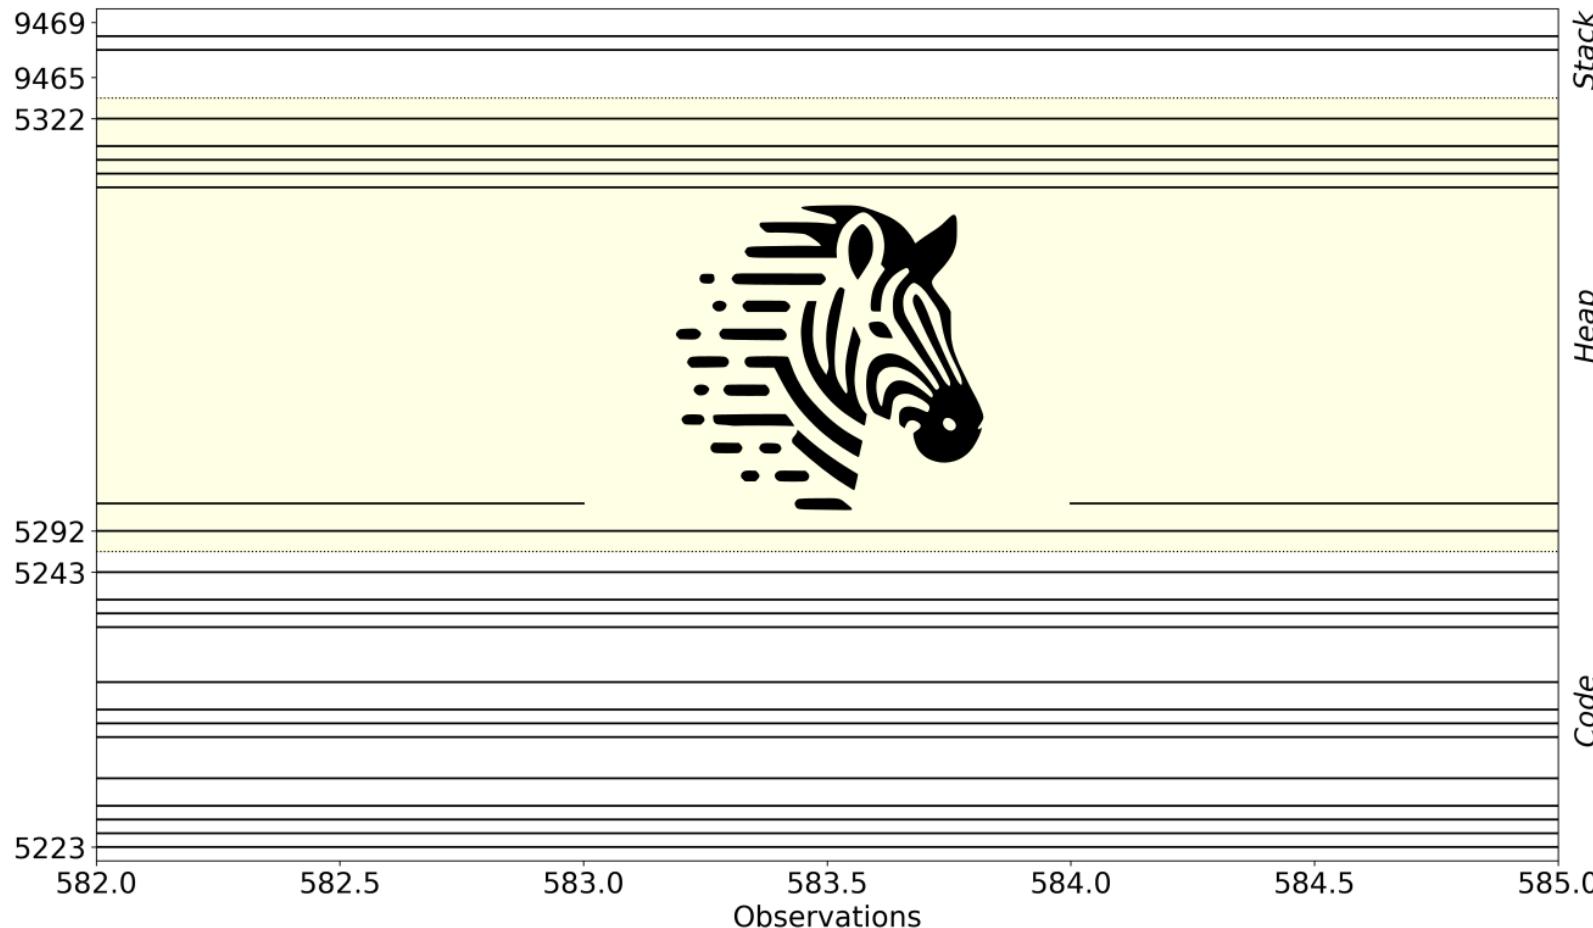
Leakage Reduction in Practice: Libjpeg TLBlur (N=10)



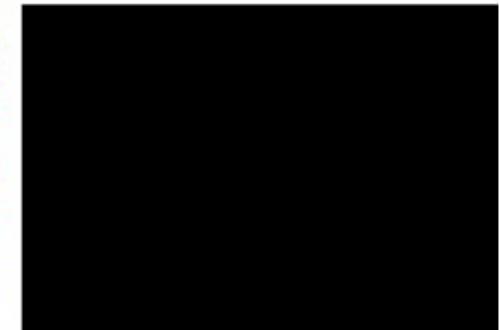
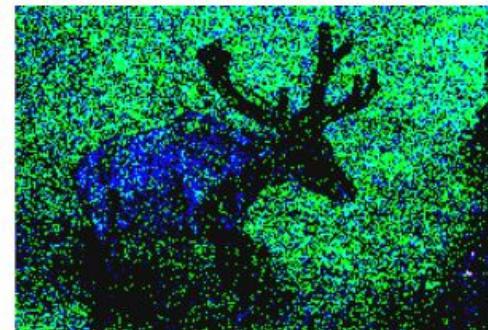
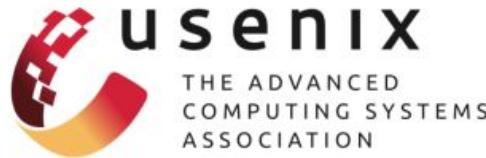
Leakage Reduction in Practice: Libjpeg TLBlur (N=20)



Leakage Reduction in Practice: Libjpeg TLBlur (N=30)



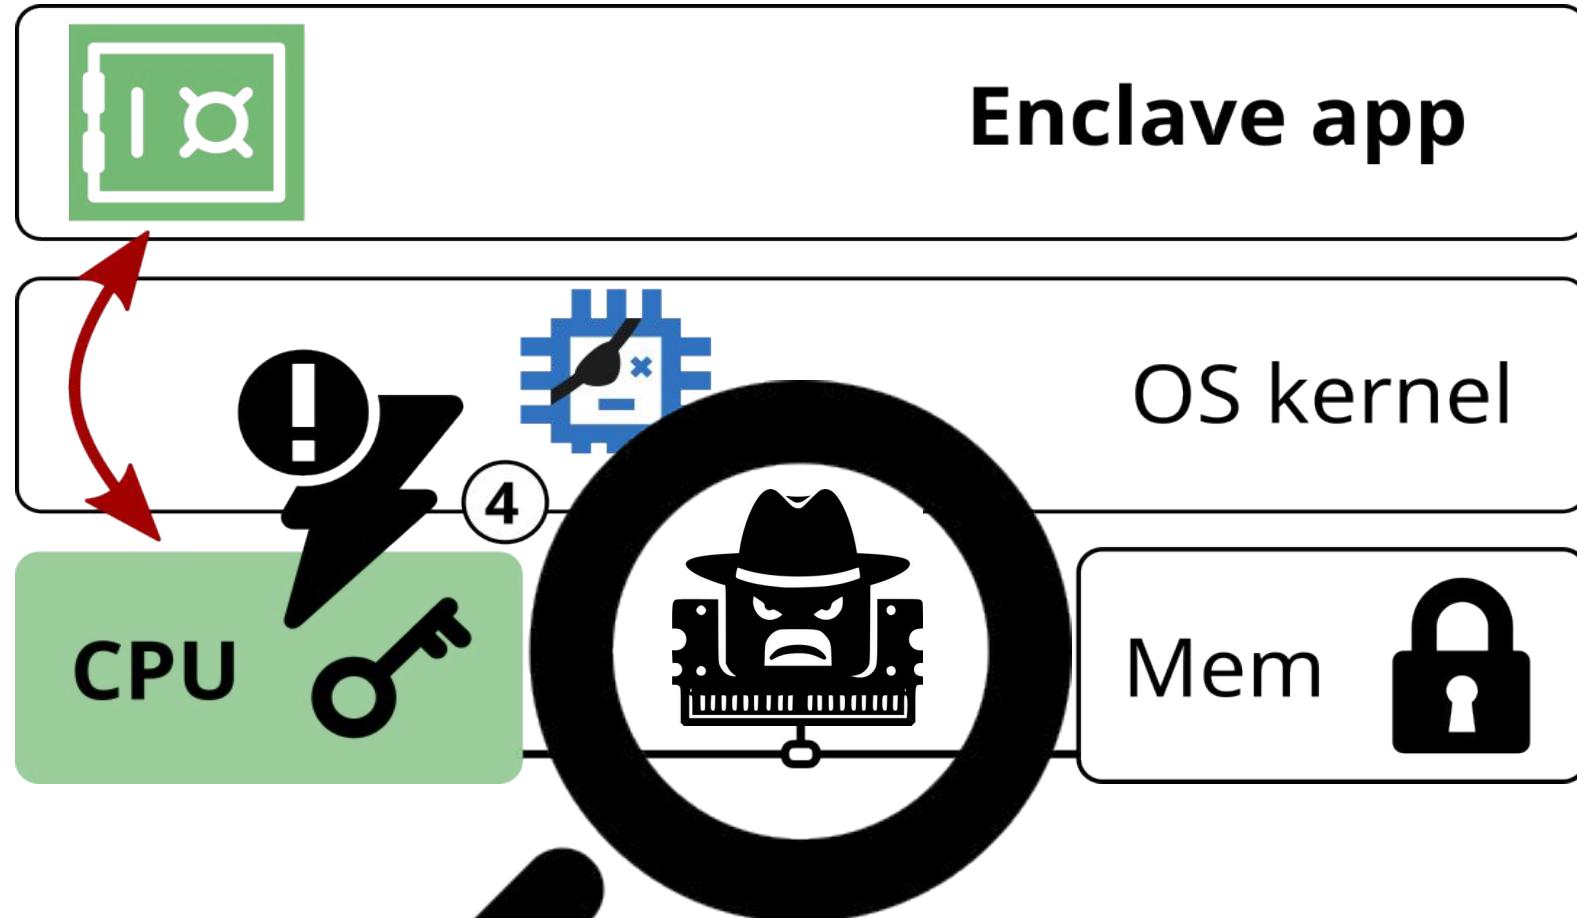
TLBlur: Compiler-Assisted Leakage Reduction in Practice



Automated “blurring” of page-access traces in space and time



Confidential-Computing: Off-Chip Attacks?



32.6. Encryption engines

In order to conceal the enclave data while it is out of the CPU package, the memory controller has an encryption engine to transparently encrypt and decrypt enclave memory.

In CPUs prior to Ice Lake, the Memory Encryption Engine (MEE) is used to encrypt pages leaving the CPU caches. MEE uses a n-ary Merkle tree with root in SRAM to maintain integrity of the encrypted data. This provides integrity and anti-replay protection but does not scale to large memory sizes because the time required to update the Merkle tree grows logarithmically in relation to the memory size.

CPUs starting from Icelake use Total Memory Encryption (TME) in the place of MEE. TME-based SGX implementations do not have an integrity Merkle tree, which means integrity and replay-attacks are not mitigated. B, it includes additional changes to prevent cipher text from being returned and SW memory aliases from being created.

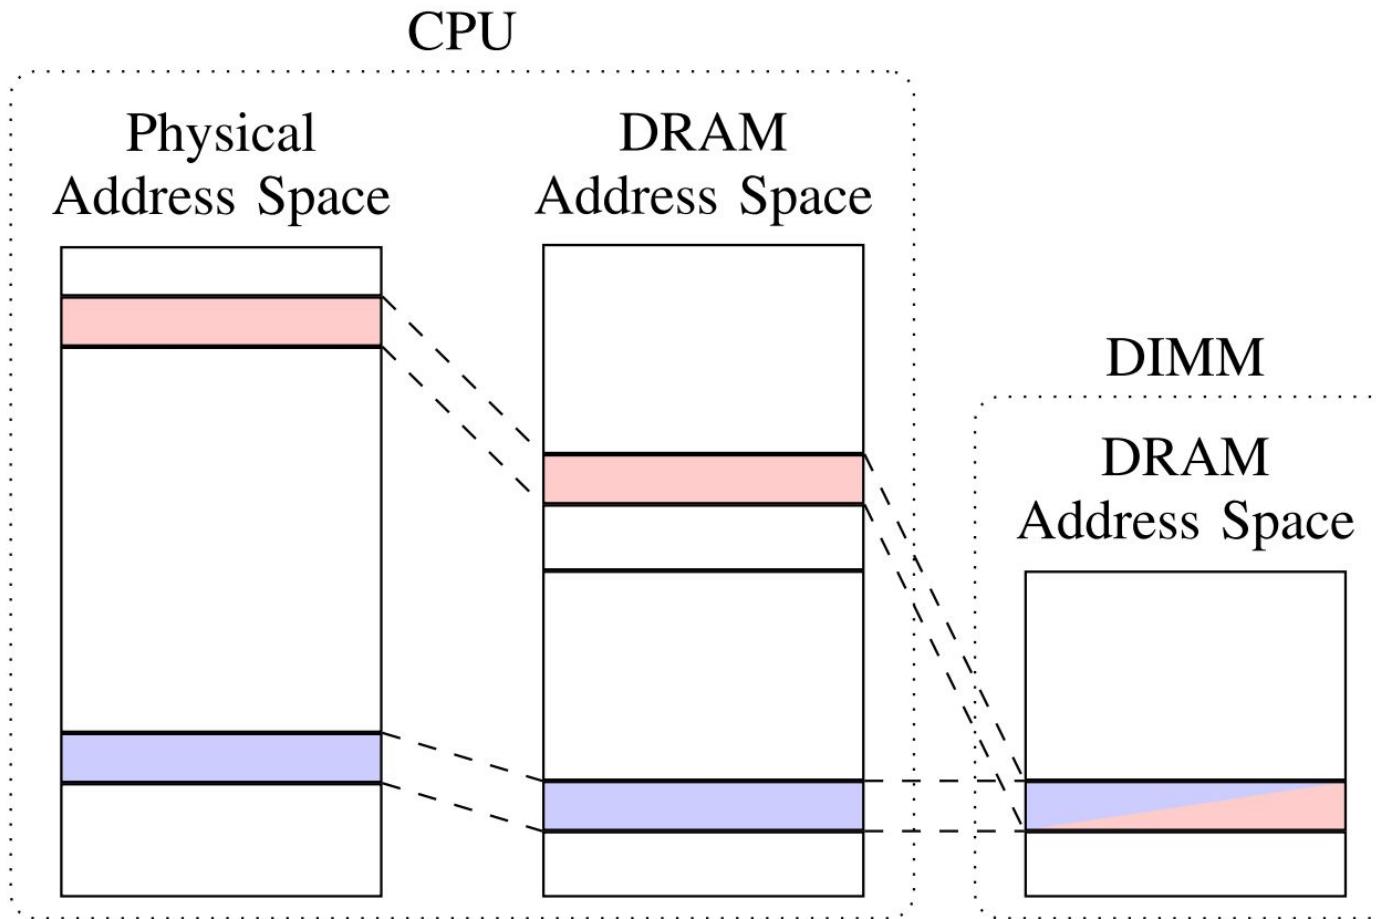
DMA to enclave memory is blocked by range registers on both MEE and TME systems (SDM section 41.10).

Background: TEE Trust in DRAM

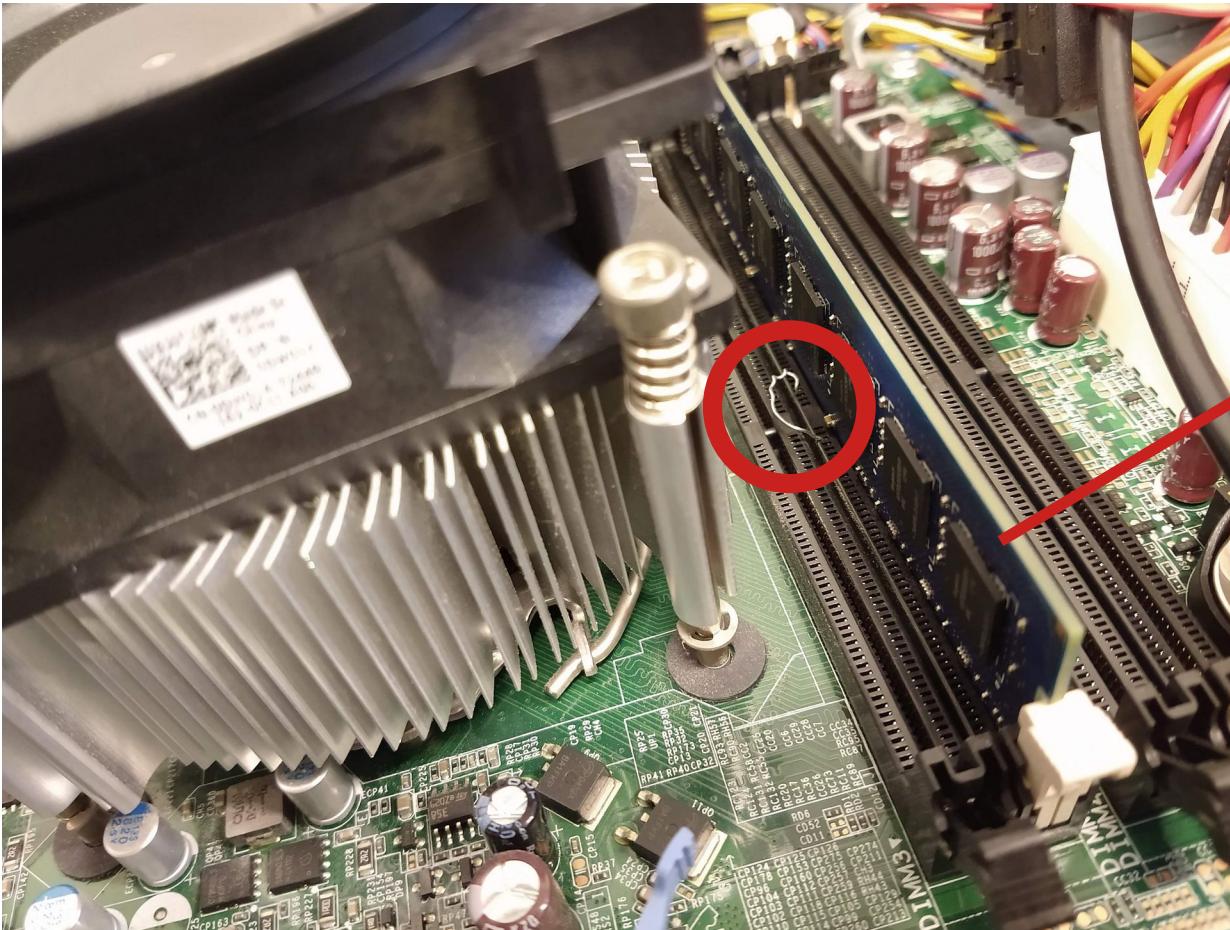
- Originally: Strong cryptographic protection
 - Limitations: Overhead, small size, ...
- Scalable solutions move away from strong cryptographic guarantees

TEE	Encryption	Guarantees		
		Confidentiality	Integrity	Freshness
Classic Intel SGX	AES-CTR	✓	✓	✓
Scalable Intel SGX	AES-XTS	✓	✗	✗
Intel TDX	AES-XTS	✓	✗	✗
AMD SEV-SNP	AES-XEX	✓	✗	✗
Arm CCA	AES-XEX/QARMA	✓	✗	✗

Idea: Memory Aliasing Attacks



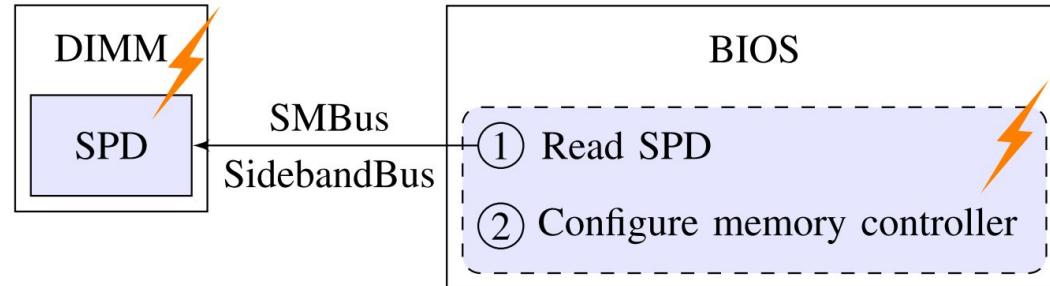
DDR3: The “Paperclip” Attack...



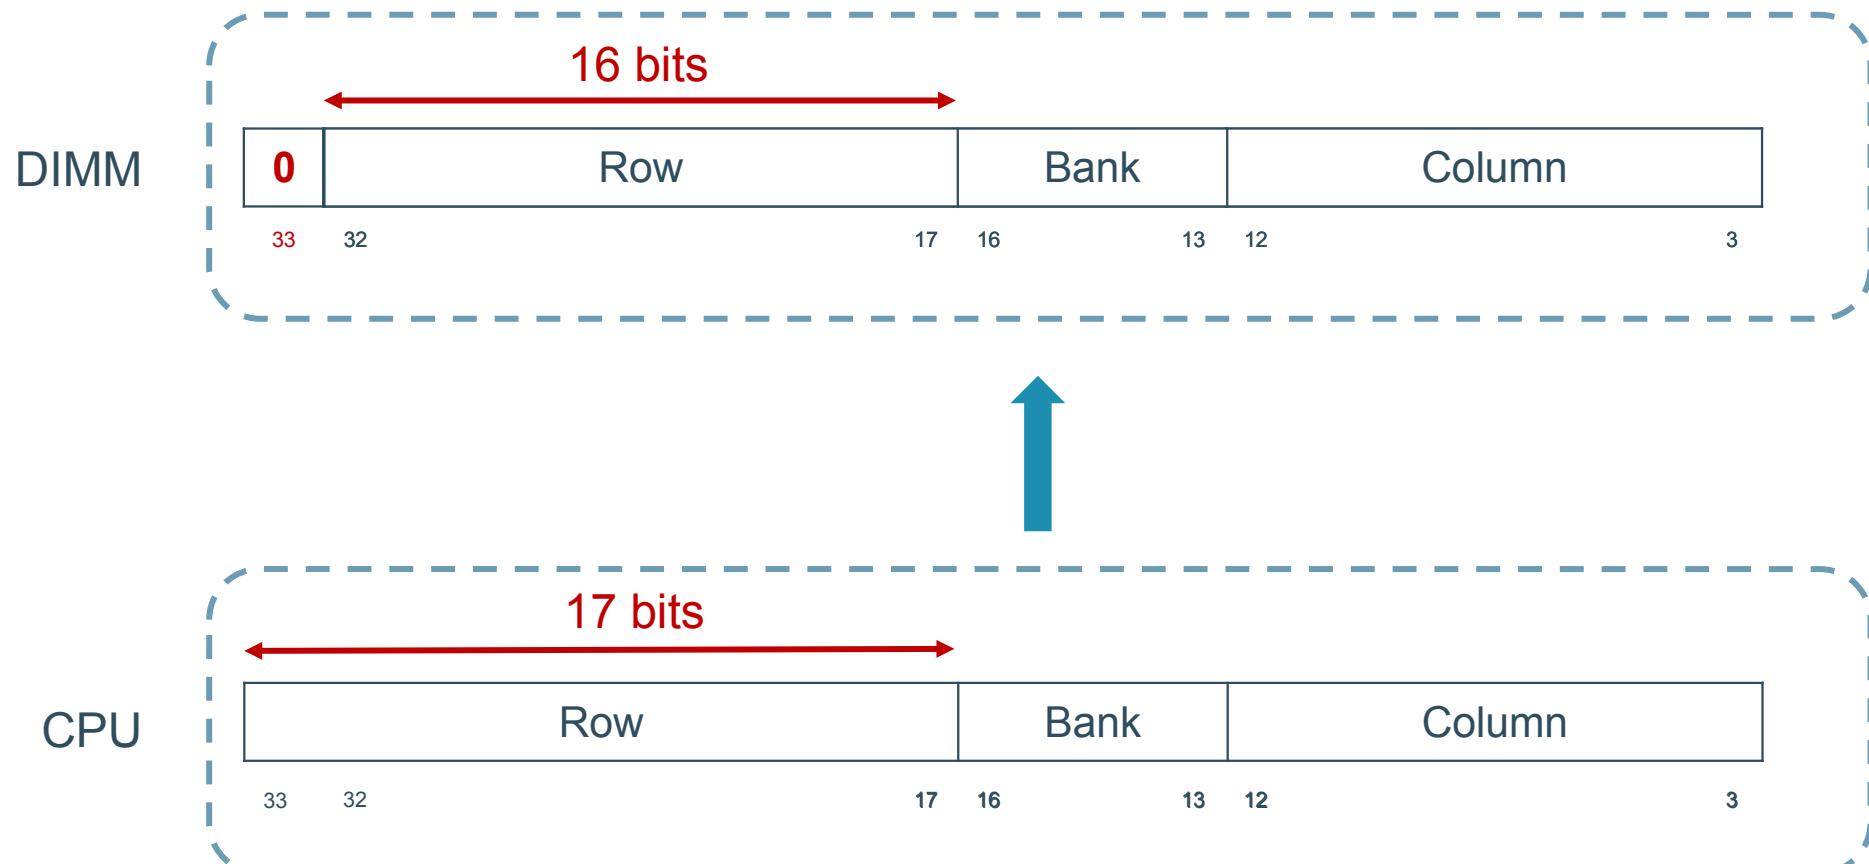
DDR3
<>
DDR4...

Background: Memory Initialization

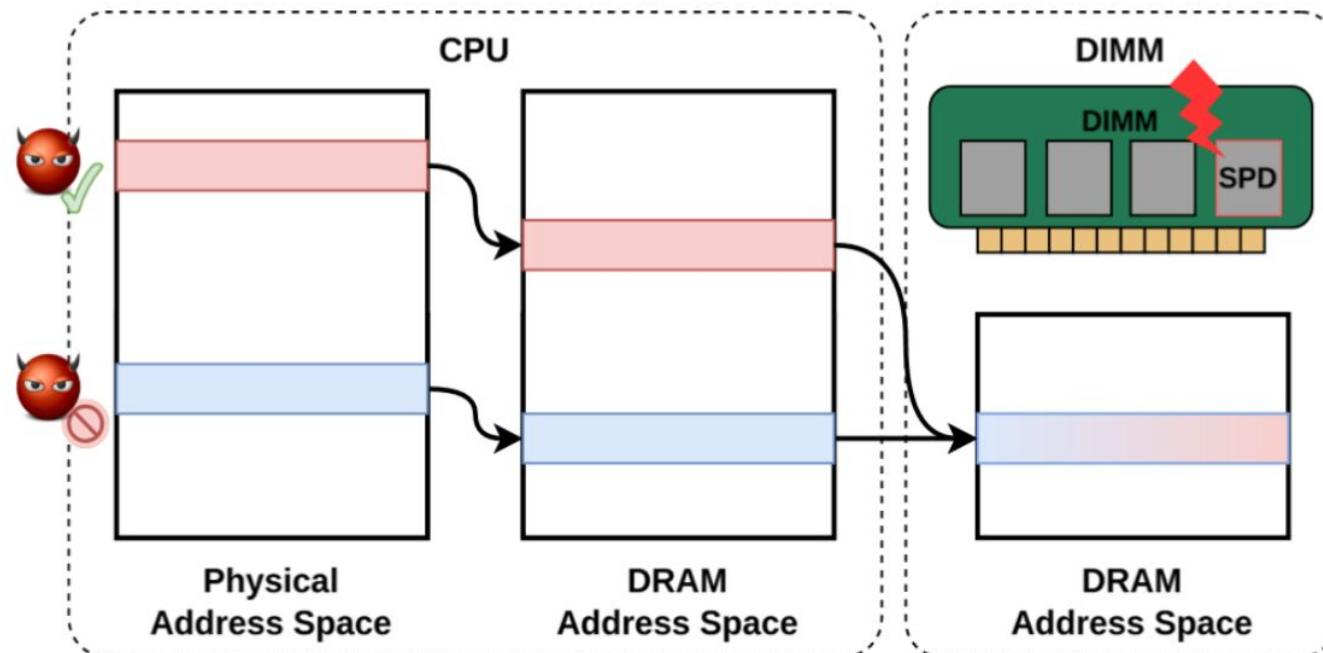
- BIOS programs memory controller based on DIMM configuration
- Incorrect configuration leads to inconsistent memory view



SPD-Based Memory Address Aliasing



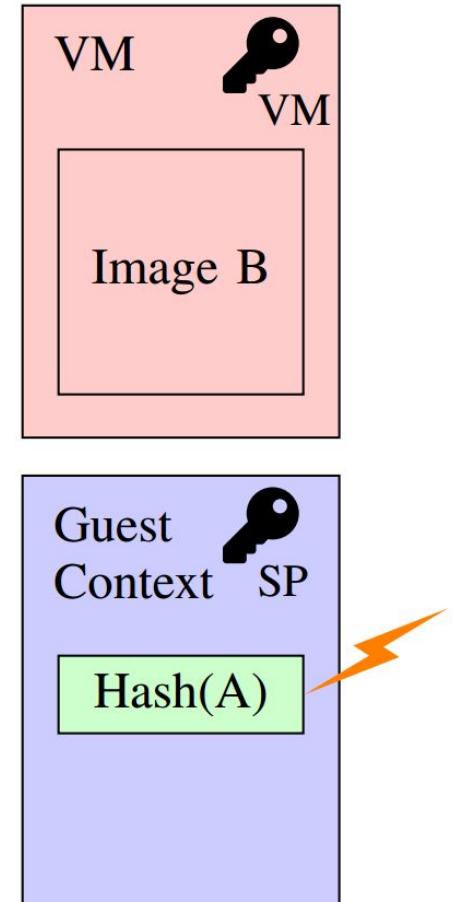
SPD-Based Memory Address Aliasing



- BIOS configures memory controller
- **Malicious SPD contents** introduces aliases

Breaking AMD SEV-SNP

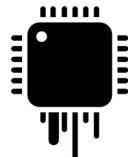
- Attestation
 - SP takes measurement of launched VM
 - Encrypted under SP key
 - No freshness, can replay launch digest



Conclusions and Take-Away



New era of **confidential computing** for the cloud and IoT



... but current architectures are **not perfect!**



Scientific understanding driven by **attacker-defender race**



Thank you!