

The Dark Side of Privilege

Understanding and Mitigating Software-Based Side Channels on Trusted Execution Environments

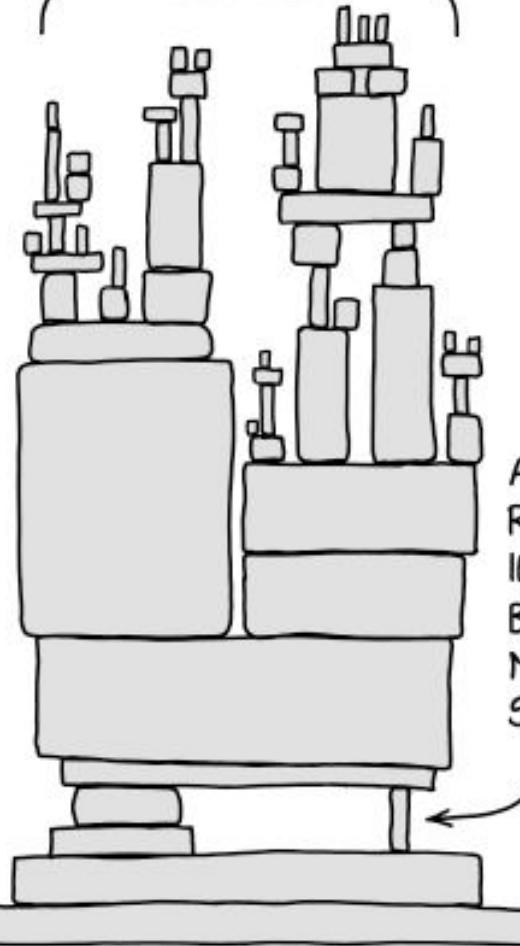
Jo Van Bulck

🏡 DistriNet, KU Leuven, Belgium 📩 jo.vanbulck@cs.kuleuven.be 🌐 vanbulck.net

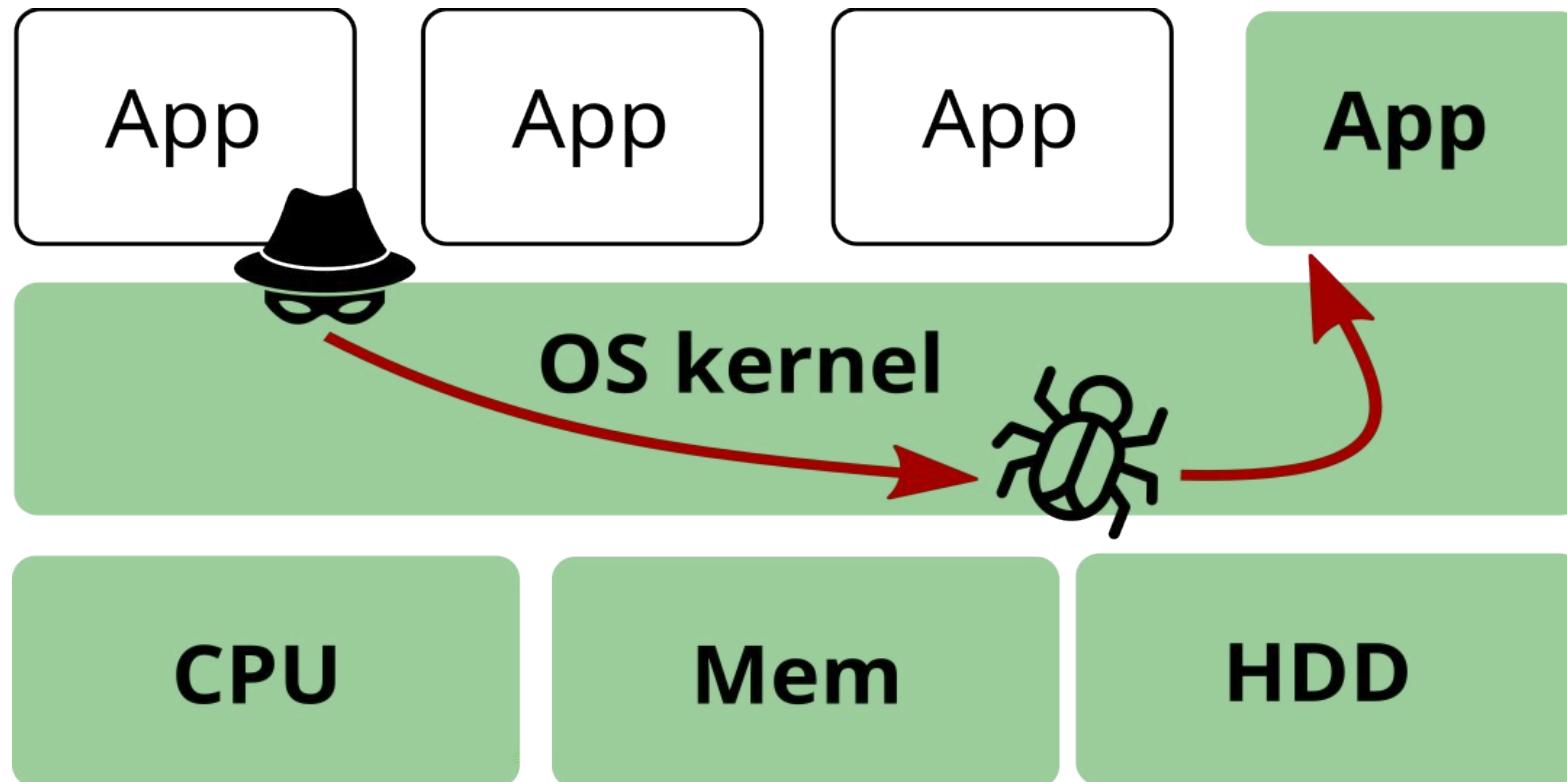
MIC-SEC Winter School, Dec 2, 2025

Trust?

ALL MODERN DIGITAL
INFRASTRUCTURE

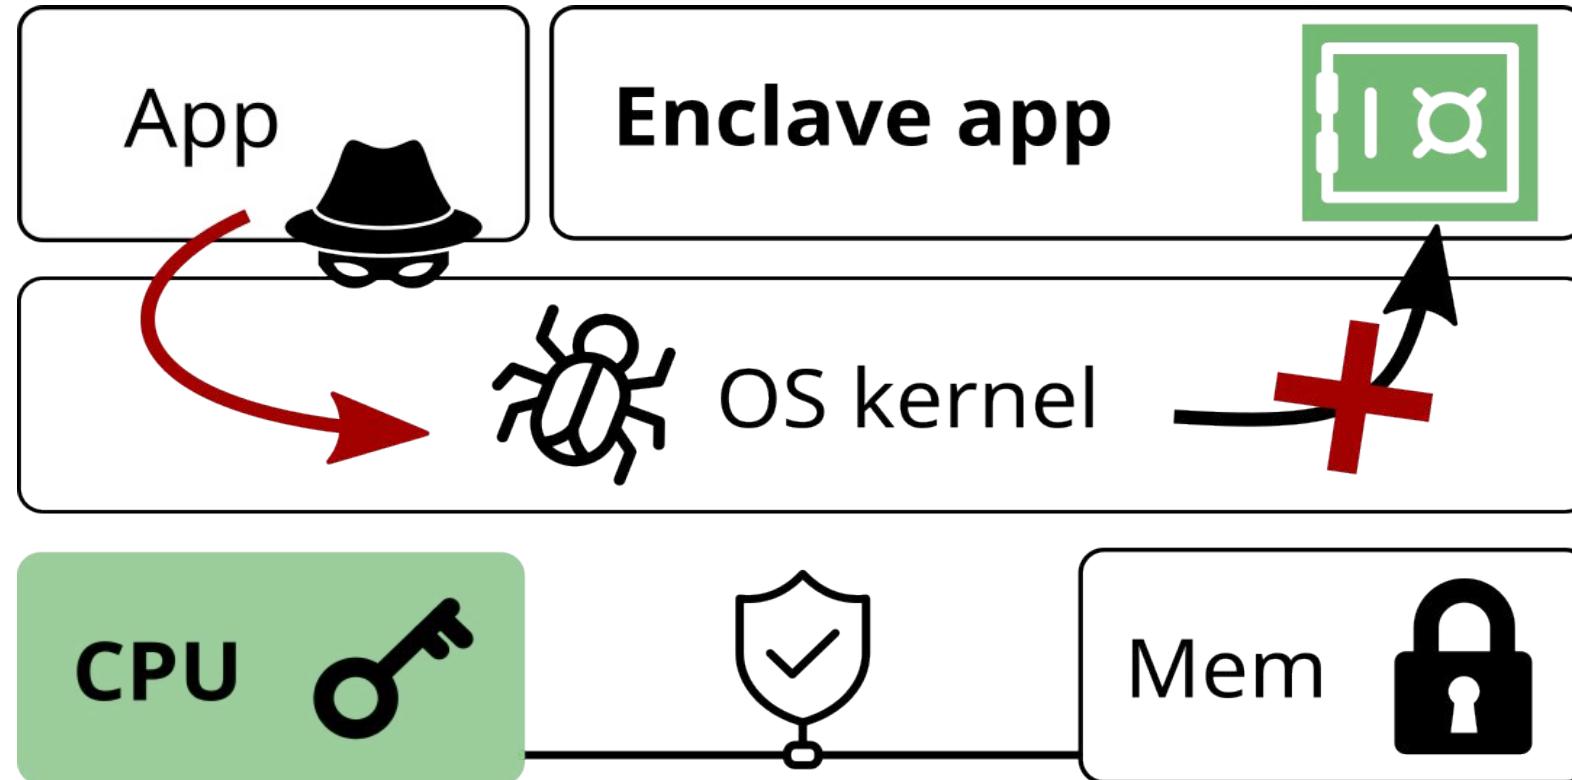


Confidential Computing: Reducing Attack Surface



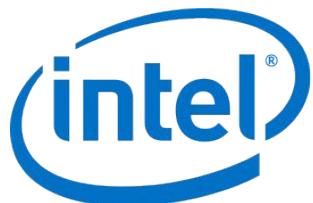
Traditional layered designs: Large **trusted computing base**

Confidential Computing: Reducing Attack Surface



Trusted execution: Hardware-level **isolation** and attestation

The Rise of Trusted Execution Environments (TEEs)




- 2004: ARM TrustZone
- 2015: **Intel Software Guard Extensions (SGX)**
- 2016: AMD Secure Encrypted Virtualization (SEV)
- 2018: IBM Protected Execution Facility (PEF)
- 2020: AMD SEV with Secure Nested Paging (SEV-SNP)
- 2022: Intel Trust Domain Extensions (TDX)
- 2023: ARM Confidential Compute Architecture (CCA)
- 2024: NVIDIA Confidential Computing



TEEs are here to stay...

Confidential Computing Isolation Paradigms



No trusted hardware


Enclave shielding

VM shielding


Intel SGX, Arm TrustZone-M

AMD SEV, Intel TDX, Arm CCA
(Arm TrustZone-A)

Confidential Computing Isolation Paradigms (Focus for Today)



“Confidential Computing Today, Just Computing Tomorrow” *



TEE Attack Research Leads the Way . . .



TEE Attack Research Leads the Way . . .



- Privileged TEE attacker models **sets the bar!**
- Idealized execution environment for attack research
- **Generalizations:** e.g., Foreshadow-NG, branch prediction, address translation, etc.

Motivation: Why Research TEE/SGX Security?



Overview About Intel News & Events Financial Info Stock Info Filings & Reports Board & Governance ESG

Overview

Press Releases

IR Calendar

Annual Stockholders' Meeting

Investor Meeting

Email Alerts

Presentations


Data Protection across the Compute Stack

Technologies such as disk- and network-traffic encryption protect data in storage and during transmission, but data can be vulnerable to interception and tampering while in use in memory. "Confidential computing" is a rapidly emerging usage category that protects data while it is in use in a Trusted Execution Environment (TEE). Intel SGX is the most researched, updated and battle-tested TEE for data center confidential computing, with the smallest attack surface within the system. It enables application isolation in private memory regions, called enclaves, to help protect up to 1 terabyte of code and data while in use.



<https://www.intc.com/news-events/press-releases/detail/1423/intel-xeon-scalable-platform-built-for-most-sensitive>


Recap: Reducing Attack Surface with Enclaves



Intel **SGX** promise: Hardware-level isolation and attestation

Not Today: Physical-Access Attacks

(See Ingrid & Jesse, Wed)



Today: Privileged Software-Based Side-Channel Attacks



Game changer: Untrusted OS → New class of powerful **side channels!**

Intel® SGX and Side-Channels

By [Simon Paul Johnson](#). Published: 03/17/2017, Last Updated: 02/27/2018

Since launching Intel® Software Guard Extensions (Intel® SGX) on 6th Generation Intel® Core™ processors in 2015, there have been a number of academic articles looking at various usage models and the security of Intel SGX. Some of these papers focus on a class of attack known as a side-channel attack, where the attacker relies on the use of a shared resource to discover information about processing occurring in some other privileged domain that it does not have direct access to.



In general, [these research papers do not demonstrate anything new or unexpected](#) about the Intel SGX architecture. Preventing side channel attacks is a matter for the enclave developer. Intel makes this clear in the security objectives for Intel SGX, which we published as part of our workshop tutorial at the International Symposium on Computer Architecture in 2015, the slides for which can be found here [\[slides 109-121\]](#), and in the [Intel® SGX SDK Developer's Manual](#).

Game Changer: Privileged “Bottom-Up” Adversary Model




Game Changer: Privileged “Bottom-Up” Adversary Model



Abuse privileged **operating system powers**
→ New and unexpected attack vectors



Privileged x86 Interface and Complexity Growth...



Systematizing the SGX Attack Landscape

	Properties						Constraints			
Attack	x86 Interface						SMT	REP	SAT	Granular
μ -arch contention	Cache priming [181, 92, 29, 225, 79]	● ○ ○ ○ ○ ○	○ ● ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	●	○	64 B
	Branch prediction [156, 59, 105]	● ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	●	○	Inst
	DRAM row buffer conflicts [263]	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	●	○	1-8 KiB
	False dependencies [180]	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	●	●	○	4 B
	Interrupt latency [256, 95, 208]	● ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	●	○	Inst
	Port contention [7]	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	●	●	○	μ -op
Control channel	Page faults [277]	● ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	○	○	4 KiB
	Page table A/D [258, 263]	● ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	○	○	4 KiB
	Page table flushing [258]	● ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	○	○	32 KiB
	Interrupt counting [182]	● ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	○	○	Inst
	IA32 segmentation faults [91]	● ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	○	○	1 B-4 KiB
	Alignment faults [254]	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	○	○	1 B
Transient	Foreshadow L1D extraction [249]	● ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	○	○	-
	Data sampling [223, 216, 211]	● ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	●	○	-
	Spectre [39, 148]	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	○	●	-
	Load value injection [251]	● ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	○	●	-
Interface	Memory safety [254, 154, 24, 266]	● ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	●	○	-
	Undervolting [188, 137, 210]	● ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	●	○	-
	Off-chip memory address bus [152]	● ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○	○	○	○	64 B




Possibly every
privileged CPU feature
can be abused!




Idea #1: Deterministic Spatial Resolution with Page Tables

Background: The Virtual Memory Abstraction




Idea: Page Faults as a Side Channel



SGX machinery protects against direct address remapping attacks


Idea: Page Faults as a Side Channel



... but untrusted address translation may **fault(!)**




Intel SGX: Page Faults as a Side Channel



□ Xu et al.: “Controlled-channel attacks: Deterministic side channels for untrusted operating systems”, Oakland 2015.


⇒ Page fault traces leak **private control data/flow**

Spatial Resolution: Page-Granular Memory Access Traces




Detailed trace of (coarse-grained) code and data accesses over time...

Spatial Resolution: Page-Granular Memory Access Traces




Spatial Resolution: Page-Granular Memory Access Traces



□ Xu et al.: “Controlled-channel attacks: Deterministic side channels for untrusted operating systems”, Oakland 2015.


... but **many faults** and a coarse-grained **4 KiB granularity**

Naive Solutions: Hiding Enclave Page Faults



- Shih et al. "T-SGX: Eradicating controlled-channel attacks against enclave programs", NDSS 2017.
- Shinde et al. "Preventing page faults from telling your secrets", AsiaCCS 2016.

Naive Solutions: Hiding Enclave Page Faults



... But stealthy attacker can learn page visits without triggering faults!

Documented Side-Effects of Address Translation


4.8 ACCESSED AND DIRTY FLAGS

For any paging-structure entry that is used during linear-address translation, bit 5 is the **accessed** flag.² For paging-structure entries that map a page (as opposed to referencing another paging structure), bit 6 is the **dirty** flag. These flags are provided for use by memory-management software to manage the transfer of pages and paging structures into and out of physical memory.

Whenever the processor uses a paging-structure entry as part of linear-address translation, it sets the accessed flag in that entry (if it is not already set).

Whenever there is a write to a linear address, the processor sets the dirty flag (if it is not already set) in the paging-structure entry that identifies the final physical address for the linear address (either a PTE or a paging-structure entry in which the PS flag is 1).

CAN'T SEE PAGE FAULTS THEY SAID




**BUT WE CAN SPY
ON PAGE TABLE MEMORY**

Telling your Secrets without Page Faults

1. Attack vector: PTE status flags:

- A(ccessed) bit
 - D(irty) bit
- ~ Also updated in enclave mode!




Telling your Secrets without Page Faults

1. Attack vector: PTE status flags:

- A(ccessed) bit
 - D(irty) bit
- ~ Also updated in enclave mode!


2. Attack vector: Unprotected page table memory:

- Cached as regular data
 - Accessed during address translation
- ~ Flush+Reload cache timing attack!



Attacking Libgcrypt EdDSA (Simplified)

```
1 if (mpi_is_secure (scalar)) {
2     /* If SCALAR is in secure memory we assume that it is the
3      secret key we use constant time operation. */
4     point_init (&tmppnt);
5
6     for (j=nbits-1; j >= 0; j--) {
7         _gcry_mpi_ec_dup_point (result, result, ctx);
8         _gcry_mpi_ec_add_points (&tmppnt, result, point, ctx);
9         point_swap_cond (result, &tmppnt, mpi_test_bit (scalar, j), ctx);
10    }
11    point_free (&tmppnt);
12 } else {
13     for (j=nbits-1; j >= 0; j--) {
14         _gcry_mpi_ec_dup_point (result, result, ctx);
15         if (mpi_test_bit (scalar, j))
16             _gcry_mpi_ec_add_points (result, result, point, ctx);
17    }
18 }
```



Attacking Libgcrypt EdDSA (Simplified)

```
1 if (mpi_is_secure (scalar)) {
2     /* If SCALAR is in secure memory we assume that it is the
3      secret key we use constant time operation. */
4     point_init (&tmppnt);
5
6     for (j=nbits-1; j >= 0; j--) {
7         _gcry_mpi_ec_dup_point (result, result, ctx);
8         _gcry_mpi_ec_add_points (&tmppnt, result, point, ctx);
9         point_swap_cond (result, &tmppnt, mpi_test_bit (scalar, j), ctx);
10    }
11    point_free (&tmppnt);
12 } else {
13     for (j=nbits-1; j >= 0; j--) {
14         _gcry_mpi_ec_dup_point (result, result, ctx);
15         if (mpi_test_bit (scalar, j))
16             _gcry_mpi_ec_add_points (result, result, point, ctx);
17    }
18 }
```

Monitor
trigger page



Memory layout

...	0x0F000
gcry_free	0xC0000
...	0xC1000
mpi_add	0xC1000
mpi_test_bit	0xC1000
...	0xC9000
mpi_ec_add_p	0xCA000
mpi_ec_mul_p	0xCA000
...	0xCA000

Attacking Libgcrypt EdDSA (Simplified)

```
1 if (mpi_is_secure (scalar)) {
2     /* If SCALAR is in secure memory we assume that it is the
3      secret key we use constant time operation. */
4     point_init (&tmppnt);
5
6     for (j=nbits-1; j >= 0; j--) {
7         _gcry_mpi_ec_dup_point (result, result, ctx);
8         _gcry_mpi_ec_add_points (&tmppnt, result, point, ctx);
9         point_swap_cond (result, &tmppnt, mpi_test_bit (scalar, j), ctx);
10    }
11    point_free (&tmppnt);
12 } else {
13     for (j=nbits-1; j >= 0; j--) {
14         _gcry_mpi_ec_dup_point (result, result, ctx);
15         if (mpi_test_bit (scalar, j))
16             _gcry_mpi_ec_add_points (result, result, point, ctx);
17    }
18 }
```

INTERRUPT




Memory layout

...	0x0F000
gcry_free	...
...	0xC0000
mpi_add	0xC1000
mpi_test_bit	...
...	0xC9000
mpi_ec_add_p	0xCA000
mpi_ec_mul_p	...

Attacking Libgcrypt EdDSA (Simplified)

```
1 if (mpi_is_secure (scalar)) {
2     /* If SCALAR is in secure memory we assume that it is the
3      secret key we use constant time operation. */
4     point_init (&tmppnt);
5
6     for (j=nbits-1; j >= 0; j--) {
7         _gcry_mpi_ec_dup_point (result, result, ctx);
8         _gcry_mpi_ec_add_points (&tmppnt, result, point, ctx);
9         point_swap_cond (result, &tmppnt, mpi_test_bit (scalar, j), ctx);
10    }
11    point_free (&tmppnt);
12 } else {
13     for (j=nbits-1; j >= 0; j--) {
14         _gcry_mpi_ec_dup_point (result, result, ctx);
15         if (mpi_test_bit (scalar, j))
16             _gcry_mpi_ec_add_points (result, result, point, ctx);
17    }
18 }
```



Attacking Libgcrypt EdDSA (Simplified)


```
1 if (mpi_is_secure (scalar)) {
2     /* If SCALAR is in secure memory we assume that it is the
3      secret key we use constant time operation. */
4     point_init (&tmppnt);
5
6     for (j=nbits-1; j >= 0; j--) {
7         _gcry_mpi_ec_dup_point (result, result, ctx);
8         _gcry_mpi_ec_add_points (&tmppnt, result, point, ctx);
9         point_swap_cond (result, &tmppnt, mpi_test_bit (scalar, j), ctx);
10    }
11    point_free (&tmppnt);
12 } else {
13     for (j=nbits-1; j >= 0; j--) {
14         _gcry_mpi_ec_dup_point (result, result, ctx);
15         if (mpi_test_bit (scalar, j))
16             _gcry_mpi_ec_add_points (result, result, point, ctx);
17    }
18 }
```



Memory layout

...	0x0F000
gcry_free	0xC0000
...	0xC1000
mpi_add	0xC1000
mpi_test_bit	0xC9000
...	0xCA000
mpi_ec_add_p	0xCA000
mpi_ec_mul_p	...


Side-channel Analysis: From Metadata Patterns to Secrets



Scientific Understanding Driven by Attacker-Defender Race...




Scientific Understanding Driven by Attacker-Defender Race...





Idea #2: Improved Temporal Resolution with Timer Interrupts



Protection from Side-Channel Attacks

Intel® SGX does not provide explicit protection from side-channel attacks. It is the enclave developer's responsibility to address side-channel attack concerns.

In general, enclave operations that require an OCall, such as thread synchronization, I/O, etc., are exposed to the untrusted domain. If using an OCall would allow an attacker to gain insight into enclave secrets, then there would be a security concern. This scenario would be classified as a side-channel attack, and it would be up to the ISV to design the enclave in a way that prevents the leaking of side-channel information.

An attacker with access to the platform can see what pages are being executed or accessed. This side-channel vulnerability can be mitigated by aligning specific code and data blocks to exist entirely within a single page.

More important, the application enclave should use an appropriate crypto implementation that is side channel attack resistant inside the enclave if side-channel attacks are a concern.

Temporal Resolution Limitations for the Page-Fault Oracle

```
1 size_t strlen (char *str)
2 {
3     char *s;
4
5     for (s = str; *s; ++s);
6     return (s - str);
7 }
```

```
1      mov   %rdi,%rax
2 1:  cmpb $0x0 ,(%rax)
3     je    2f
4     inc   %rax
5     jmp   1b
6 2:  sub   %rdi,%rax
7     retq
```


⇒ tight loop: 4 instructions, single memory operand, single code + data page

Counting strlen loop iterations?



Note: Page-fault attacks cannot make progress for 1 code + data page

Temporal Resolution Limitations for the Page-Fault Oracle



Counting `strlen` loop iterations?




Progress requires both pages present (non-faulting) \leftrightarrow page fault oracle

Building the Side-Channel Oracle with Execution Timing?



Too noisy: modern x86 processors are lightning fast...



Challenge: Side-Channel Sampling Rate



Slow
shutter speed




Medium
shutter speed



Fast
shutter speed

SGX-Step: Executing Enclaves one Instruction at a Time



SGX-Step: Executing Enclaves one Instruction at a Time



SGX-Step



<https://github.com/jovanbulck/sgx-step>

Watch

22

Star

245


Fork

52



Van Bulck et al., "SGX-Step: A Practical Attack Framework for Precise Enclave Execution Control", SysTEX 2017.

SGX-Step: Executing Enclaves one Instruction at a Time



SGX-Step Demo: Single-Stepping Password Comparison


```
jo@breuer:~/sgx-step-demo$ sudo ./app ■
```

SGX-Step: Enabling a New Line of High-Resolution Attacks

Yr	Venue	Paper	Step	Use Case	Drv
'15	S&P	Ctrl channel	~ Page	Probe (page fault)	✓
'16	ESORICS	AsyncShock	~ Page	Exploit (mem safety)	-
'17	CHES	CacheZoom	X >1	Probe (L1 cache)	✓
'17	ATC	Hahnel et al.	X 0 - >1	Probe (L1 cache)	✓
'17	USENIX	BranchShadow	X 5 - 50	Probe (BPU)	X
'17	USENIX	Stealthy PTE	~ Page	Probe (page table)	✓
'17	USENIX	DarkROP	~ Page	Exploit (mem safety)	✓
'17	SysTEX	SGX-Step	✓ 0 - 1	Framework	✓
'18	ESSoS	Off-limits	✓ 0 - 1	Probe (segmentation)	✓
'18	AsiaCCS	Single-trace RSA	~ Page	Probe (page fault)	✓
'18	USENIX	Foreshadow	✓ 0 - 1	Probe (transient exec)	✓
'18	EuroS&P	SgxPectre	~ Page	Exploit (transient)	✓
'18	CHES	CacheQuote	X >1	Probe (L1 cache)	✓
'18	ICCD	SGXlinger	X >1	Probe (IRQ latency)	X
'18	CCS	Nemesis	✓ 1	Probe (IRQ latency)	✓
'19	USENIX	Spoiler	✓ 1	Probe (IRQ latency)	✓
'19	CCS	ZombieLoad	✓ 0 - 1	Probe (transient exec)	✓
'19	CCS	Fallout	-	Probe (transient exec)	✓
'19	CCS	Tale of 2 worlds	✓ 1	Exploit (mem safety)	✓
'19	ISCA	MicroScope	~ 0 - Page	Framework	X
'20	CHES	Bluethunder	✓ 1	Probe (BPU)	✓
'20	USENIX	Big troubles	~ Page	Probe (page fault)	✓
'20	S&P	Plundervolt	-	Exploit (undervolt)	✓
'20	CHES	Viral primitive	✓ 1	Probe (IRQ count)	✓
'20	USENIX	CopyCat	✓ 1	Probe (IRQ count)	✓
'20	S&P	LVI	✓ 1	Exploit (transient)	✓

Yr	Venue	Paper	Step	Use Case	Drv
'20	CHES	A to Z	~ Page	Probe (page fault)	✓
'20	CCS	Déjà Vu NSS	~ Page	Probe (page fault)	✓
'20	MICRO	PTHammer	-	Probe (page walk)	✓
'21	USENIX	Frontal	✓ 1	Probe (IRQ latency)	✓
'21	S&P	CrossTalk	✓ 1	Probe (transient exec)	✓
'21	CHES	Online template	✓ 1	Probe (IRQ count)	✓
'21	NDSS	SpeechMiner	-	Framework	✓
'21	S&P	Platypus	✓ 0 - 1	Probe (voltage)	✓
'21	DIMVA	Aion	✓ 1	Probe (cache)	✓
'21	CCS	SmashEx	✓ 1	Exploit (mem safety)	✓
'21	CCS	Util::Lookup	✓ 1	Probe (L3 cache)	✓
'22	USENIX	Rapid prototyping	✓ 1	Framework	✓
'22	CT-RSA	Kalyna expansion	✓ 1	Probe (L3 cache)	✓
'22	SEED	Enclzyer	-	Framework	✓
'22	NordSec	Self-monitoring	~ Page	Defense (detect)	✓
'22	AutoSec	Robotic vehicles	✓ 1 - >1	Exploit (timestamp)	✓
'22	ACSAC	MoLE	✓ 1	Defense (randomize)	✓
'22	USENIX	AEPIC	✓ 1	Probe (I/O device)	✓
'22	arXiv	Confidential code	✓ 1	Probe (IRQ latency)	✓
'23	ComSec	FaultMorse	~ Page	Probe (page fault)	✓
'23	CHES	HQC timing	✓ 1	Probe (L3 cache)	✓
'23	ISCA	Belong to us	✓ 1	Probe (BPU)	✓
'23	USENIX	BunnyHop	✓ 1	Probe (BPU)	✓
'23	USENIX	DownFall	✓ 0 - 1	Probe (transient exec)	✓
'23	USENIX	AEX-Notify	✓ 1	Defense (prefetch)	✓

SGX-Step: A Versatile Open-Source Attack Framework



Page-table manipulation






Interrupt latency




SGX-Step



High-resolution probing




Interrupt counting



Zero-step replaying


Example: Deterministic Instruction Counting Attacks



We can *count instructions!*

Inputs to `expf()` from 10^{-2} to 10^2 fall into into **11 different count classes** (8 unique).


Example: Augmented Instruction-Level Page-Access Traces



Nemesis: Extracting IRQ Latency Traces with SGX-Step




Enclave x-ray: IRQ latency leaks instruction-level μ -arch timing!



Nemesis: Extracting IRQ Latency Traces with SGX-Step




Enclave x-ray: Spotting **high-latency** instructions



Nemesis: Extracting IRQ Latency Traces with SGX-Step




Enclave x-ray: Zooming in on `bsearch` function




De-Anonymizing Enclave Lookups with Interrupt Latency

Binary search: Find 40 in {20, 30, 40, 50, 80, 90, 100}




De-Anonymizing Enclave Lookups with Interrupt Latency

Adversary: Infer secret lookup in known array




De-Anonymizing Enclave Lookups with Interrupt Latency

Goal: Infer lookup → reconstruct bsearch control flow




De-Anonymizing Enclave Lookups with Interrupt Latency

Goal: Infer lookup → reconstruct bsearch control flow





De-Anonymizing Enclave Lookups with Interrupt Latency


⇒ Sample **instruction latencies** in secret-dependent path



Interrupt-Latency Attacks Beyond Intel SGX




 **TEXAS
INSTRUMENTS**




The Intel logo, featuring the word "intel" in a lowercase sans-serif font inside a blue swoosh.

Single-Stepping Beyond Intel SGX




**THE CODE DOESN'T WORK...
WHY?**



**THE CODE WORKS...
WHY?**

Root-Causing SGX-Step: Aiming the Timer Interrupt



Root-Causing SGX-Step: Microcode Assists to the Rescue!

PTE A-bit	Mean (cycles)	Stddev (cycles)
A=1	27	30
A=0	666	55



3. Assisted PT walk



1. Clear PTE A-bit




2. TLB flush



Arm timer


ERESUME

NOP₁




12

Root-Causing SGX-Step: Microcode Assists to the Rescue!






AEX-Notify: Hardware-Software Co-Design Solution



AEX-Notify: Hardware-Software Co-Design Solution

We implemented a fast, constant-time decoder (CTD)



CHAPTER 8

ASYNCHRONOUS ENCLAVE EXIT NOTIFY AND THE EDECCSSA USER LEAF FUNCTION

8.1 INTRODUCTION

Asynchronous Enclave Exit Notify (AEX-Notify) is an extension to Intel® SGX that allows Intel SGX enclaves to be notified after an asynchronous enclave exit (AEX) has occurred. EDECCSSA is a new Intel SGX user leaf function (ENCLU[EDECCSSA]) that can facilitate AEX notification handling, as well as software exception handling. This chapter provides information about changes to the Intel SGX architecture that support AEX-Notify and ENCLU[EDECCSSA].

The following list summarizes the a
details are provided in Section 8.3)

- SECS.ATTRIBUTES.AEXNOTIFY
- TCS.FLAGS.AEXNOTIFY: This en
- SSA.GPRSGX.AEXNOTIFY: Enclave-writable byte that allows enclave software to dynamically enable/disable AEX notifications.



SGX-Step led to new x86 processor instructions!

→ shipped in millions of devices \geq 4th Gen Xeon CPU

An AEX notification is delivered by ENCLU[ERESUME] when the following conditions are met:

Intel AEX Notify Support Prepped For Linux To Help Enhance SGX Enclave Security

Written by Michael Larabel in Intel on 6 November 2022 at 06:01 AM EST. [5 Comments](#)



Future Intel CPUs and some existing processors via a microcode update will support a new feature called the Asynchronous EXit (AEX) notification mechanism to help with Software Guard Extensions (SGX) enclave security. Patches for the Linux kernel are pending for implementing this Intel AEX Notify support with capable processors.

Intel's Asynchronous EXit (AEX) notification mechanism lets SGX enclaves run a handler after an AEX event. Those handlers can be used for things like mitigating SGX-Step as an attack framework for precise enclave execution control.



Code 1 in intel/linux-sgx X

intel sdk/trts/linux/trts_mitigation.S

```
48 * Description:  
49 *      The file provides mitigations for SGX-Step  
50 */  
51 * Function:  
52 *      constant_time_apply_sgxstep_mitigation_and_continue_execution  
53 *      Mitigate SGX-Step and return to the point at which the  
most recent  
54 *      interrupt/exception occurred.
```





SGX-Step led to changes in major OSs and enclave SDKs


There's a Catch...

Finally note that our proposed mitigation does not protect against interrupting enclaves and observing application code and data page accesses at a coarse-grained 4 KiB spatial resolution. In contrast to the fine-grained, instruction-granular interrupt-driven attacks we consider in this work, such controlled-channel attacks have received ample attention [18, 47, 56, 59] from the research community.


Why Mitigating Single-Stepping is Not Enough




Original (left), Xu et al. (middle), our attack with AEX-Notify single-stepping defense (right)




Libjpeg: AEX-Notify's Temporal Reduction in Practice




Libjpeg: AEX-Notify's Temporal Reduction in Practice




Libjpeg: AEX-Notify's Temporal Reduction in Practice




Idea: TLB as a “Filter” to Hide Page Accesses




TLBlur: Self-Monitoring and Restoring Enclave Page Accesses




Instrumentation to Self-Monitor Page Accesses at Runtime




Leakage Reduction in Practice: Libjpeg Single-Stepping




Leakage Reduction in Practice: Libjpeg Page Faults




Leakage Reduction in Practice: Libjpeg TLBlur (N=10)






Leakage Reduction in Practice: Libjpeg TLBlur (N=20)



Leakage Reduction in Practice: Libjpeg TLBlur (N=30)



TLBlur: Compiler-Assisted Leakage Reduction in Practice



Automated “blurring” of page-access traces in space and time






Idea #3: Model-Specific Registers and Performance-Monitoring Counters?

Software-Based DVFS Attacks



Software-Based DVFS Attacks



□ Lipp et al. "PLATYPUS: Software-based Power Side-Channel Attacks on x86", S&P 21.



Requirements: Software-only attacker → *Metadata + direct data leakage(!)*


Software-Based DVFS Attacks



□ Murdock et al. "Plundervolt: Software-Based Fault Injection Attacks against Intel SGX", S&P 20.



Requirements: Software-only attacker → Fault injection (*integrity* breach)



How a little bit of undervolting can cause a lot of problems



Warning

Altering clock frequency or voltage may:

- damage or reduce the useful life of the processor and other system components
- may reduce system stability and performance.

Product warranties may not apply if the processor is operated beyond its specifications. Check with the manufacturers of system and components for additional details.



I agree

I agree, don't show again

Cancel

Intel Power Management




Fig. 1. Layout of the undocumented undervolting MSR with address 0x150.

Plundervolt: Will it Fault?

```
uint64_t multiplier = 0x1122334455667788;
uint64_t var = 0xdeadbeef * multiplier;

while (var == 0xdeadbeef * multiplier)
{
    var = 0xdeadbeef;
    var *= multiplier;
}
var ^= 0xdeadbeef * multiplier;
```

Plundervolt: Will it Fault?




Fig. 3. Base voltage (blue) and voltage for first fault (orange) vs. CPU frequency for the i3-7100U-A

Plundervolt: Differential Fault Analysis of AES-NI in SGX

```
[Enclave] plaintext: 5ABB97CCFE5081A4598A90E1CEF1BC39
[Enclave] CT1: DE49E9284A625F72DB87B4A559E814C4 <- faulty
[Enclave] CT2: BDFADCE3333976AD53BB1D718DFC4D5A <- correct
```

input to round 10:

```
[Enclave] 1: CD58F457 A9F61565 2880132E 14C32401
[Enclave] 2: AEEBC19C D0AD3CBA A0BCBAFA C0D77D9F
```

input to round 9:

```
[Enclave] 1: 6F6356F9 26F8071F 9D90C6B2 E6884534
[Enclave] 2: 6F6356C7 26F8D01F 9DF7C6B2 A4884534
```

input to round 8:

```
[Enclave] 1: 1C274B5B 2DFD8544 1D8AEAC0 643E70A1
[Enclave] 2: 1C274B5B 2DFD8544 1D8AEAC0 646670A1
```

Plundervolt: Beyond Crypto—Inducing Memory-Safety Faults





Figure 4. The address of element $a[i]$ in an array is computed as $\&a[0] + i * \text{sizeof}(\text{elem_t})$.

Plundervolt: Beyond Crypto—Inducing Memory-Safety Faults






Conclusions and Take-Away



New era of **confidential computing** for the cloud and IoT



... but current architectures are **not perfect!**



Scientific understanding driven by **attacker-defender race**



Thank you!